

Transformer

Assoc. Prof. Peerapon Vateekul, Ph.D.

Department of Computer Engineering,
Faculty of Engineering, Chulalongkorn University

Peerapon.v@chula.ac.th

(Credit to TA's & This slide is based on 2110572: Natural Language Processing Systems)

Attention mechanism

Published as a conference paper at ICLR 2015

NEURAL MACHINE TRANSLATION BY JOINTLY LEARNING TO ALIGN AND TRANSLATE

Dzmitry Bahdanau
Jacobs University Bremen, Germany

KyungHyun Cho **Yoshua Bengio***
Université de Montréal

ABSTRACT

Neural machine translation is a recently proposed approach to machine translation. Unlike the traditional statistical machine translation, the neural machine translation aims at building a single neural network that can be jointly tuned to maximize the translation performance. The models proposed recently for neural machine translation often belong to a family of encoder-decoders and encode a source sentence into a fixed-length vector from which a decoder generates a translation. In this paper, we conjecture that the use of a fixed-length vector is a bottleneck in improving the performance of this basic encoder-decoder architecture, and propose to extend this by allowing a model to automatically (soft-)search for parts of a source sentence that are relevant to predicting a target word, without having to form these parts as a hard segment explicitly. With this new approach, we achieve a translation performance comparable to the existing state-of-the-art phrase-based system on the task of English-to-French translation. Furthermore, qualitative analysis reveals that the (soft-)alignments found by the model agree well with our intuition.

1 INTRODUCTION

Attention Is All You Need

Ashish Vaswani*
Google Brain
avaswani@google.com

Noam Shazeer*
Google Brain
noam@google.com

Niki Parmar*
Google Research
nikip@google.com

Jakob Uszkoreit*
Google Research
usz@google.com

Llion Jones*
Google Research
llion@google.com

Aidan N. Gomez* †
University of Toronto
aidan@cs.toronto.edu

Lukasz Kaiser*
Google Brain
lukaszkaiser@google.com

Illia Polosukhin* ‡
illia.polosukhin@gmail.com

Attention Is All You Need. 31st Conference on Neural Information Processing Systems (NIPS 2017), Long Beach, CA, USA.

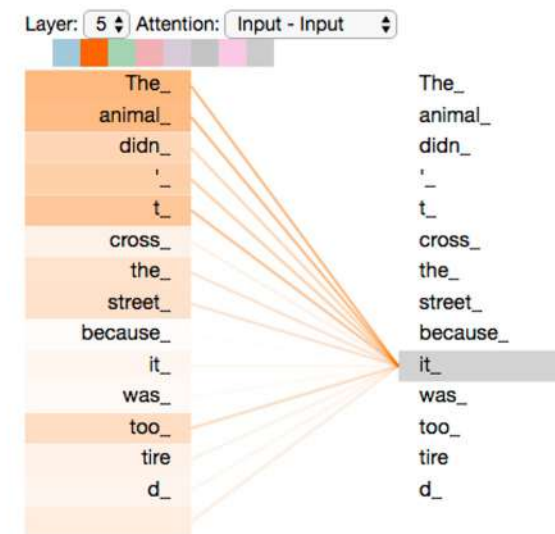
Reference: Bahdanau, Dzmitry, Kyunghyun Cho, and Yoshua Bengio. "Neural machine translation by jointly learning to align and translate." ICLR(2015).

+ Scaled Dot-Product Attention (1)

- Introduced in Attention is all you need (Vaswani et al., 2017)
- **NO** recurrence nor convolution
- Widely used today in all Transformer-based model
- “Relating different positions of a single sequence in order to compute a representation of the sequence”

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

where the query, keys, values, and output are all vectors
and d_k is a number.



+ Scaled Dot-Product Attention (2)

4

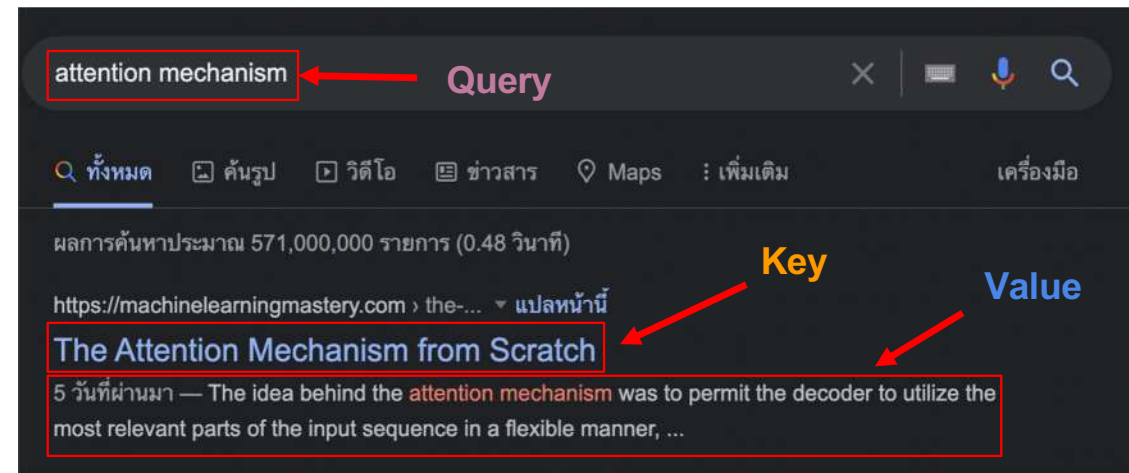
What are the “query”, “key”, and “value” vectors?


As an analogy, think of Google Search.

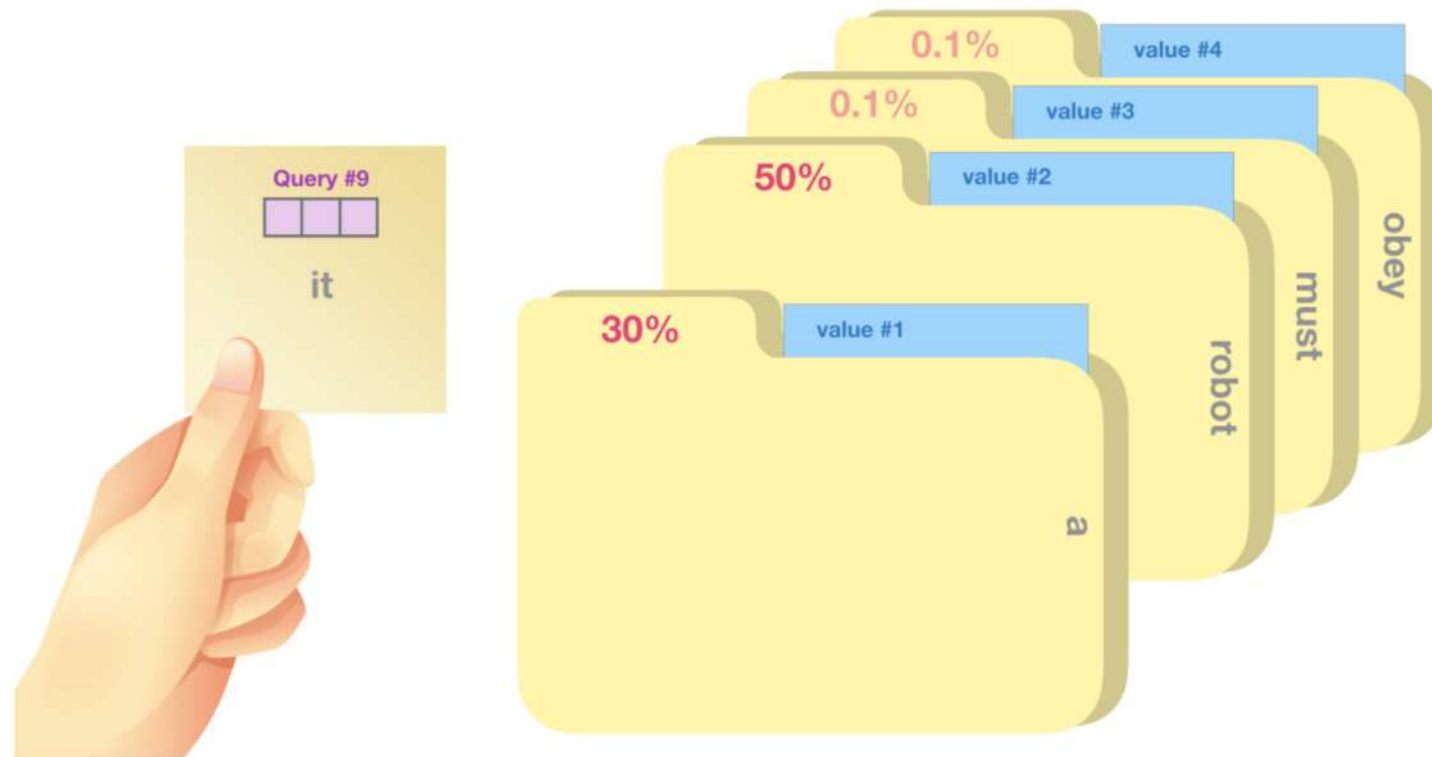
Query - what we want to know

Key - how to index information

Value - what kind of info is in each website







Input1
Input2


Thinking
Machines

Embedding


X

W^Q


Q



\times




$=$




X

W^K


K



\times




$=$




X

W^V



V



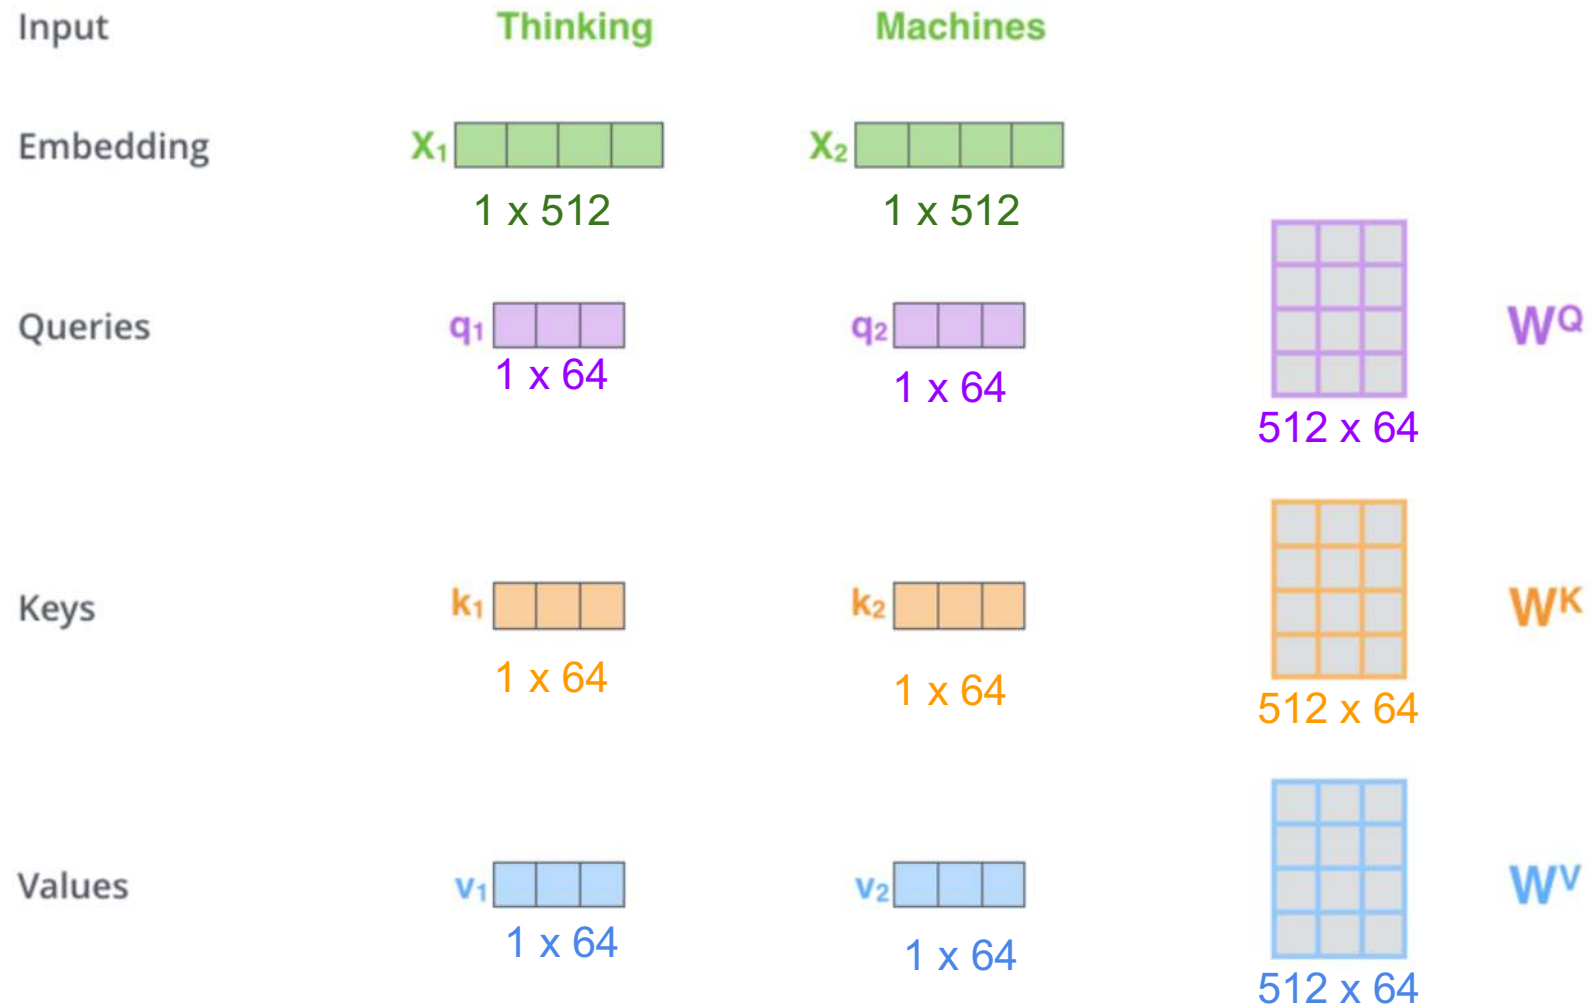
\times



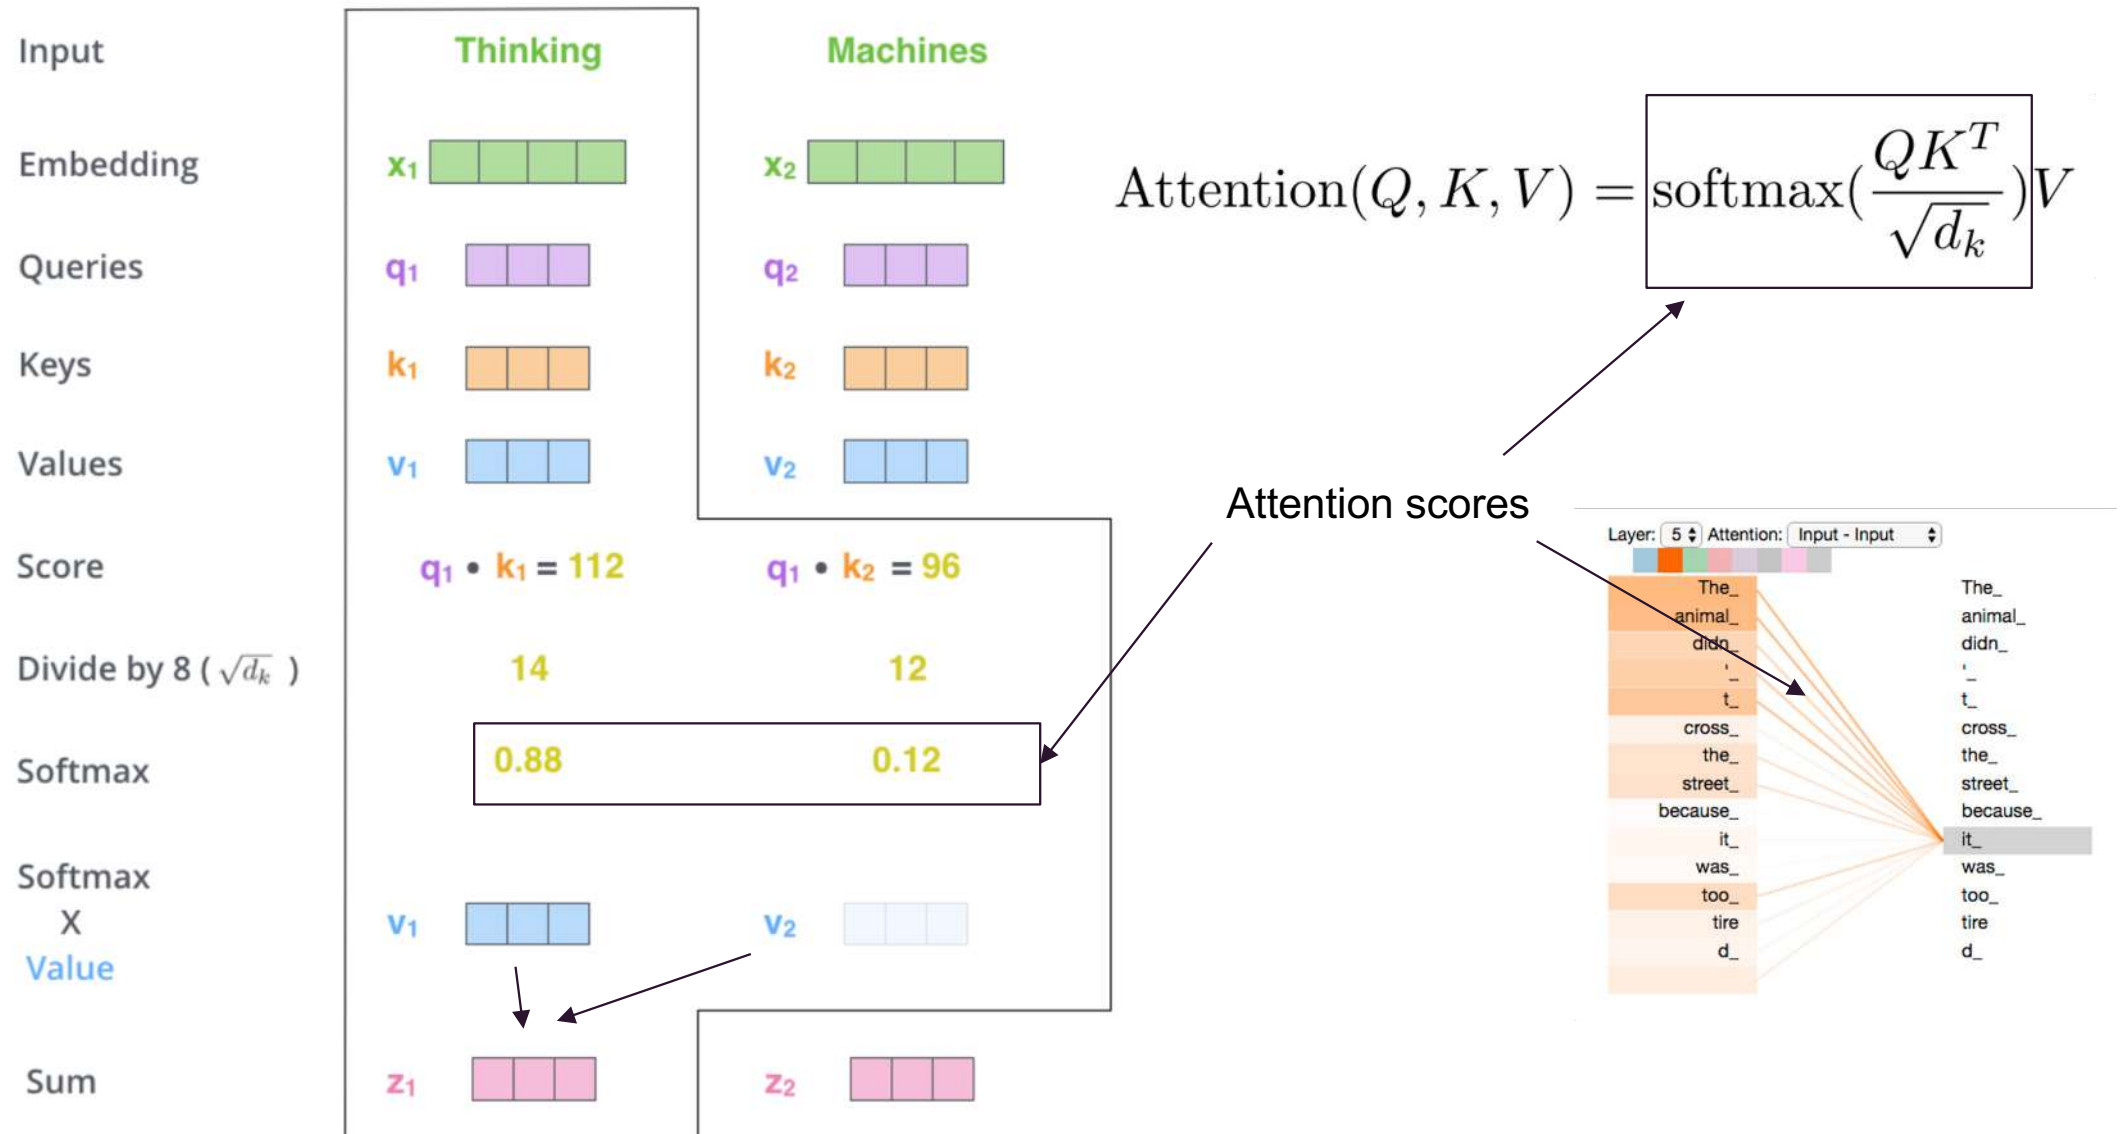
$=$




W^Q , W^K , and W^V are linear layers




Multiplying x_1 by the W^Q weight matrix produces q_1 , the "query" vector associated with that word. We end up creating a "query", a "key", and a "value" projection of each word in the input sentence.



$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$



Scaled Dot-Product Attention



Attention Is All You Need

Ashish Vaswani*
Google Brain
avaswani@google.com

Noam Shazeer*
Google Brain
noam@google.com

Niki Parmar*
Google Research
nikip@google.com

Jakob Uszkoreit*
Google Research
usz@google.com

Llion Jones*
Google Research
llion@google.com

Aidan N. Gomez* †
University of Toronto
aidan@cs.toronto.edu

Łukasz Kaiser*
Google Brain
lukaszkaizer@google.com



Illia Polosukhin* ‡
illia.polosukhin@gmail.com

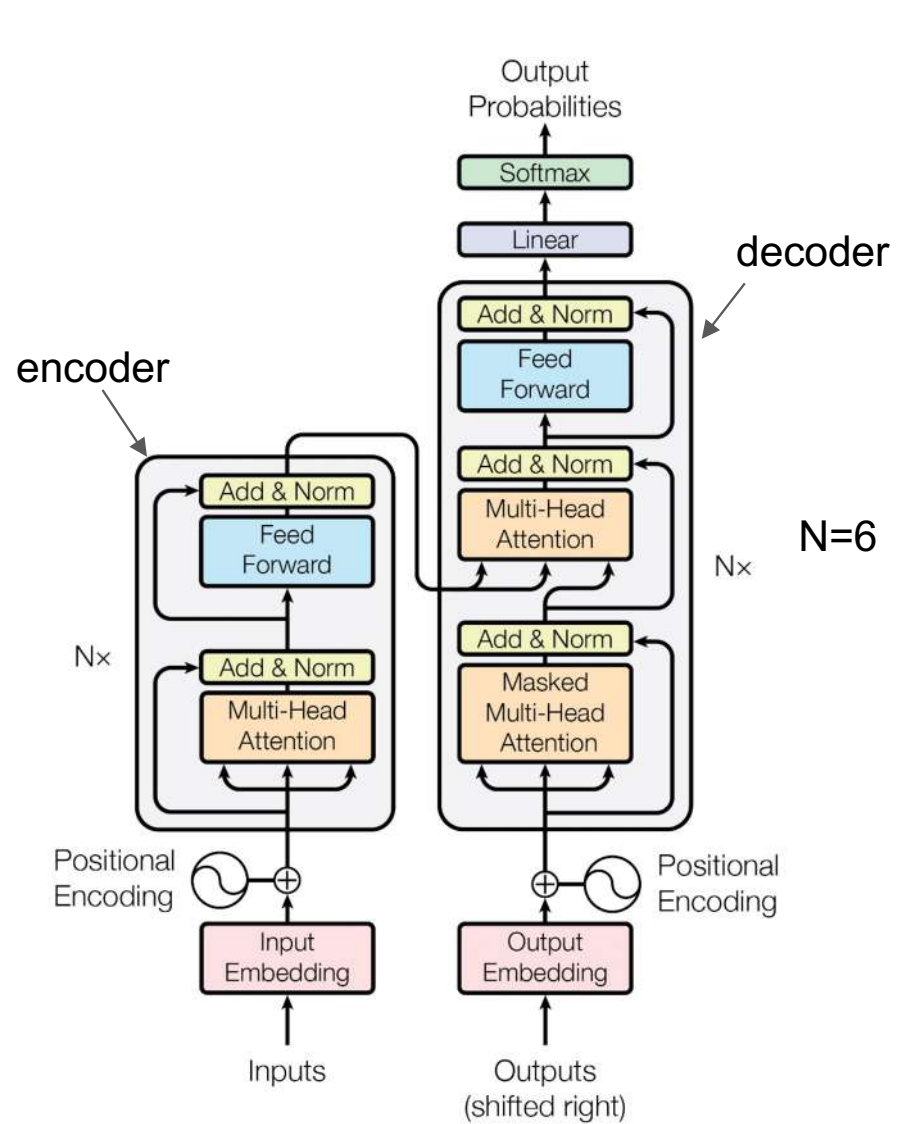
6 Dec 2017

The Transformer

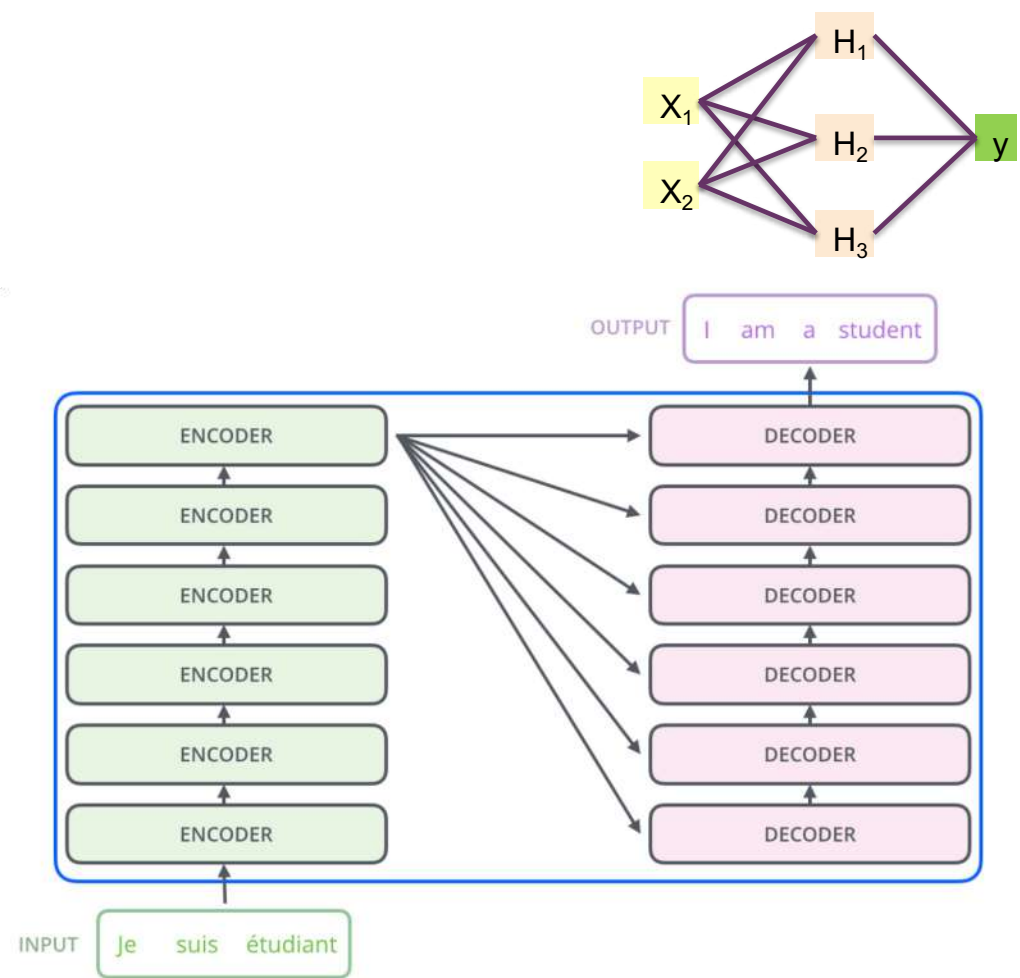
Only rely on the attention mechanism - **No RNN!**

- More parallelizable -> Faster training time
- Better at longer sequence

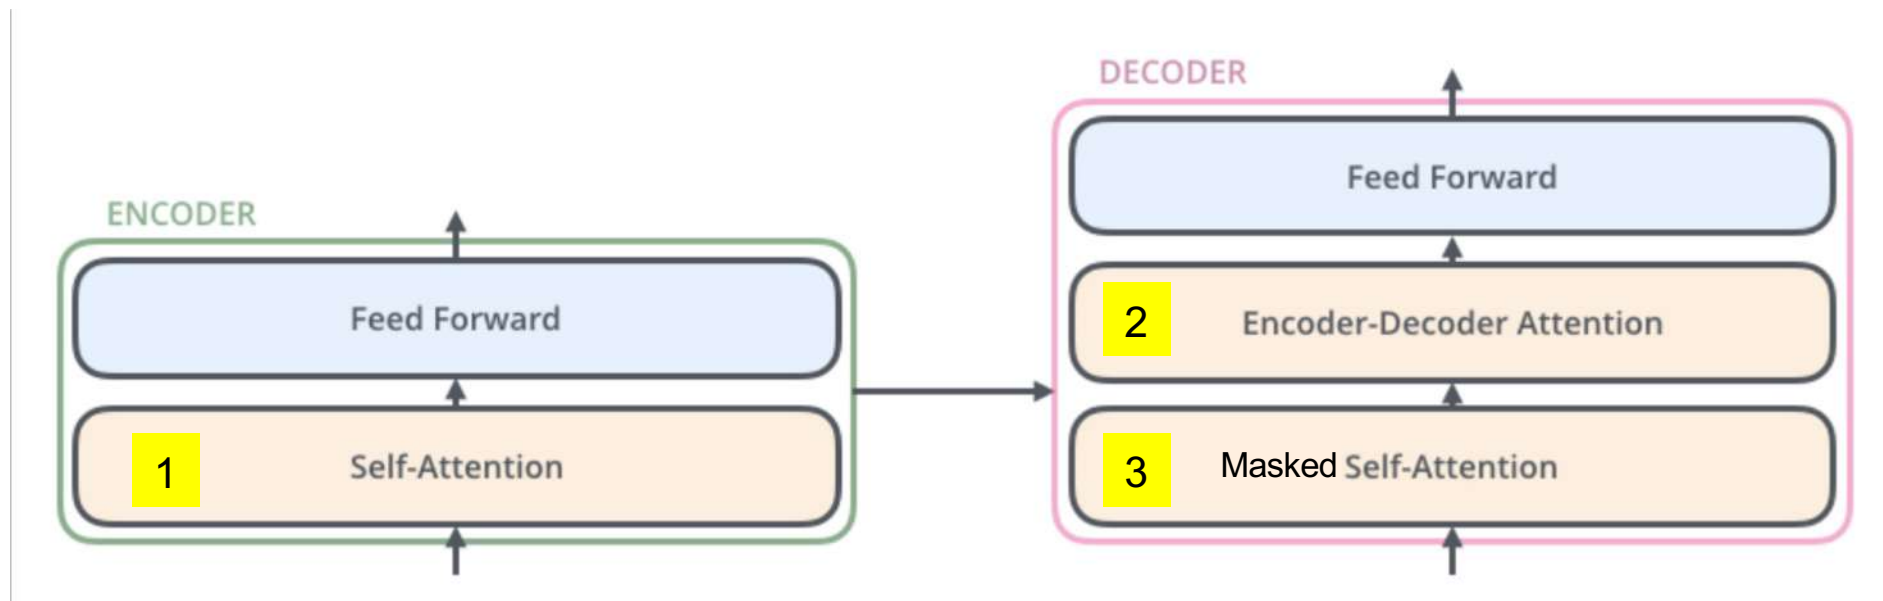




=



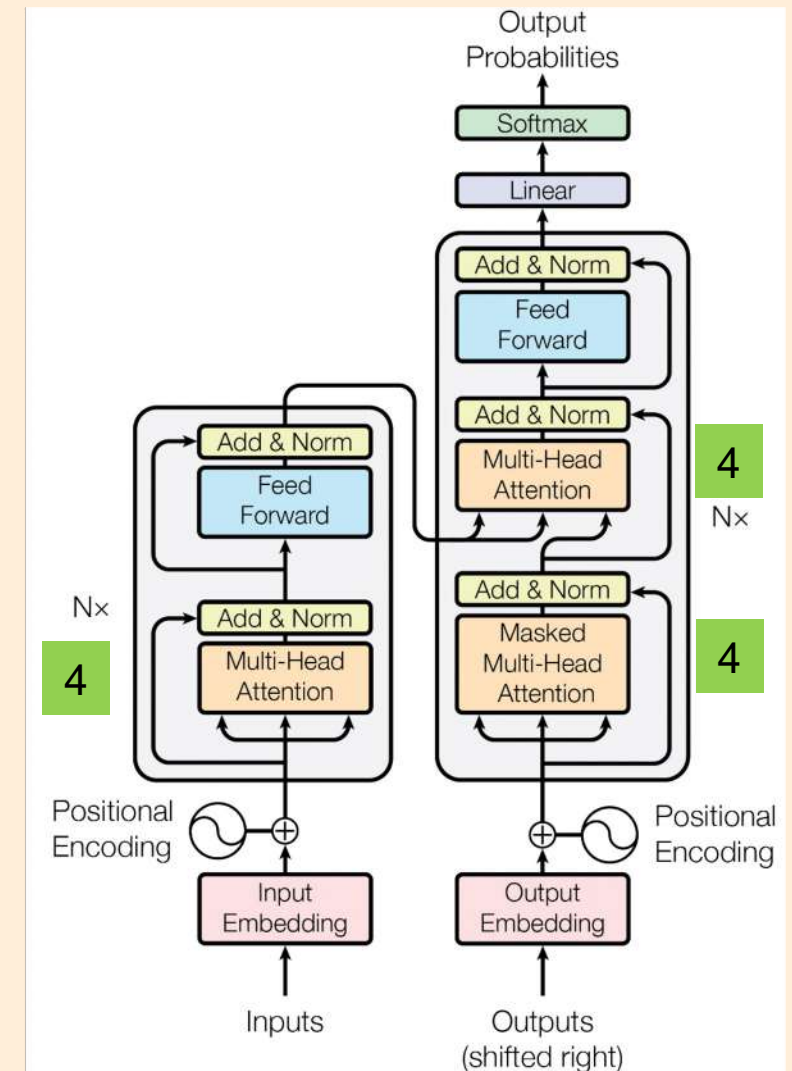
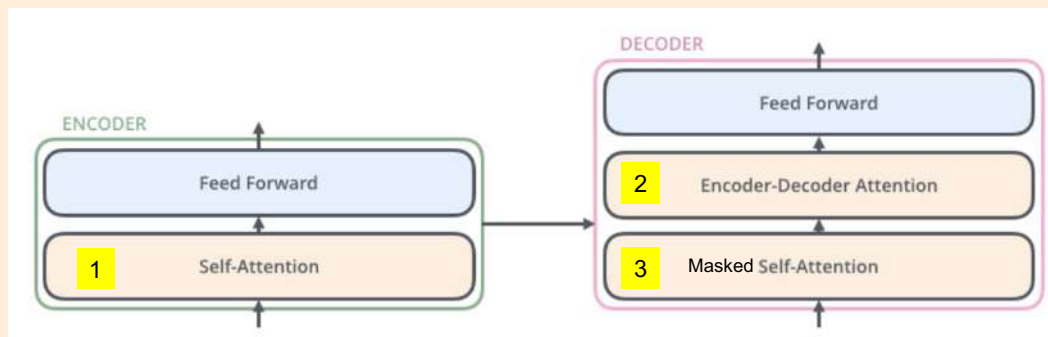
In each layer..



Scaled Dot-Product Attention

There are many variations used in the Transformer:

1. Self Attention
2. Encoder-decoder Attention
3. Masked Self Attention
4. Multi-headed Self Attention





1) Self Attention

Q, K, and V are all from the same input.

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

where the query, keys, values, and output are all vectors and d_k is a number.

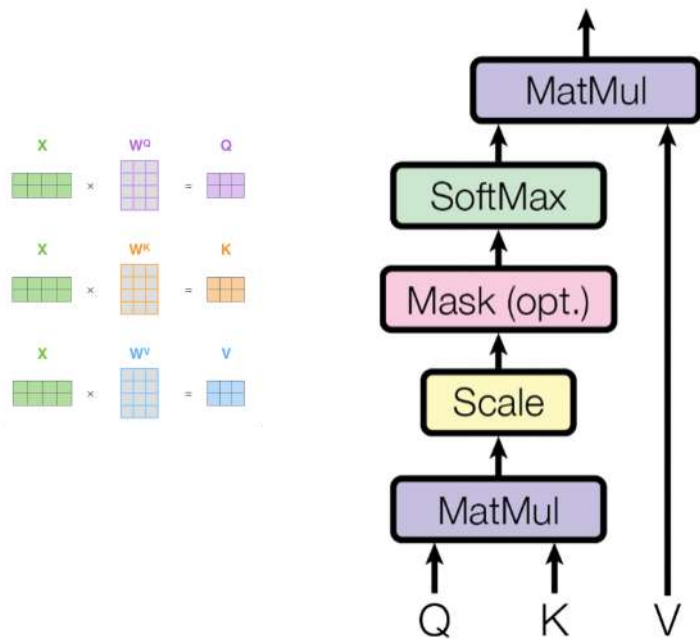
Scaled Dot-Product Attention



My name is Sam → ฉัน ชื่อ แซม

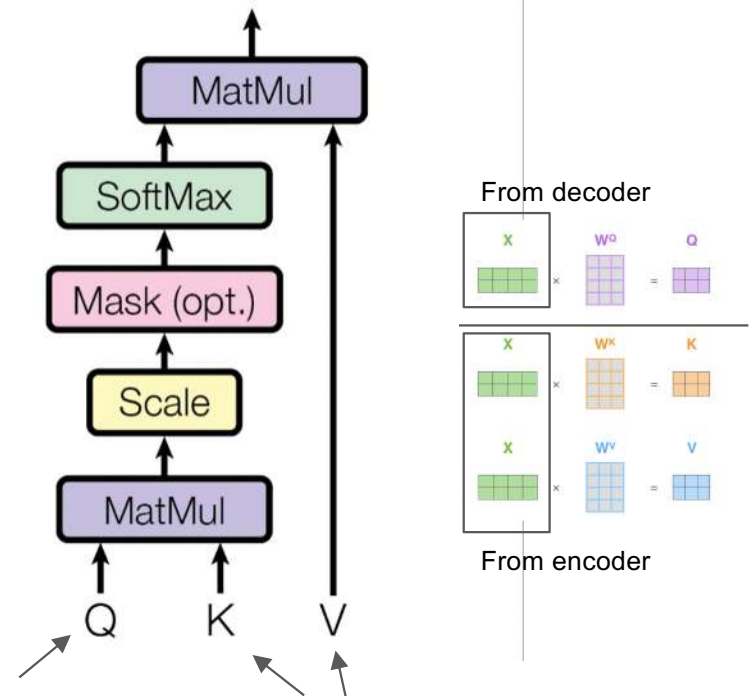
2) Encoder-Decoder Attention

Self Attention



ALL from same input

Encoder-Decoder Attention
(or Cross Attention)



From decoder



From encoder

My name is Sam → ฉันชื่อแซม

3) Masked Self Attention

Prevent the model from **seeing into the future** to preserve the autoregressive property.

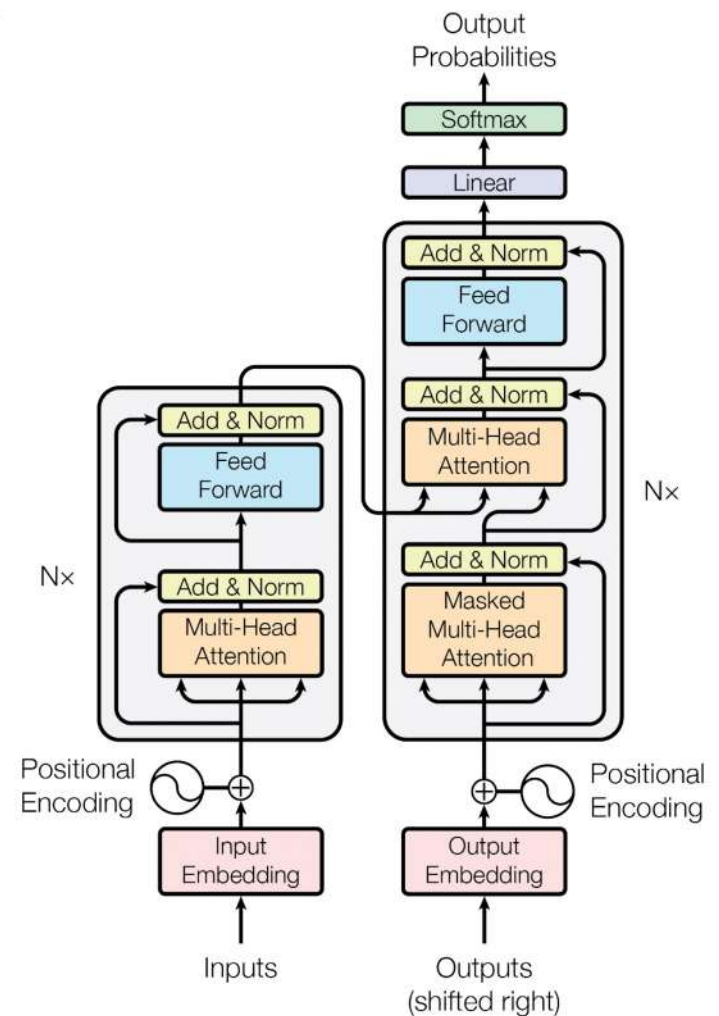

Used in *decoder* only (so the model can't see the answer).



Masked Self Attention

Prevent the model from **seeing into the future** to preserve the autoregressive property.

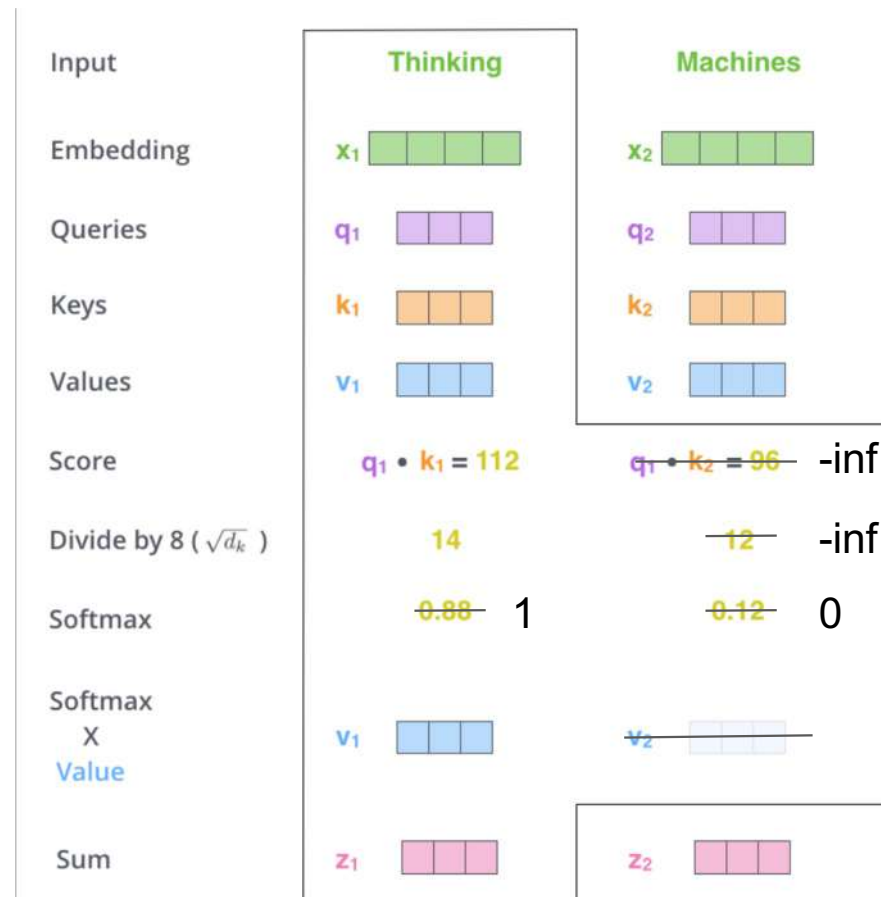
Used in *decoder* only (so the model can't see the answer).



I like dogs <eos>

<bos> ฉัน ชอบ สุนัข

Masked Self Attention






4) Multi-head Self Attention

Multi-head attention allows the model to jointly attend to information from different representations.

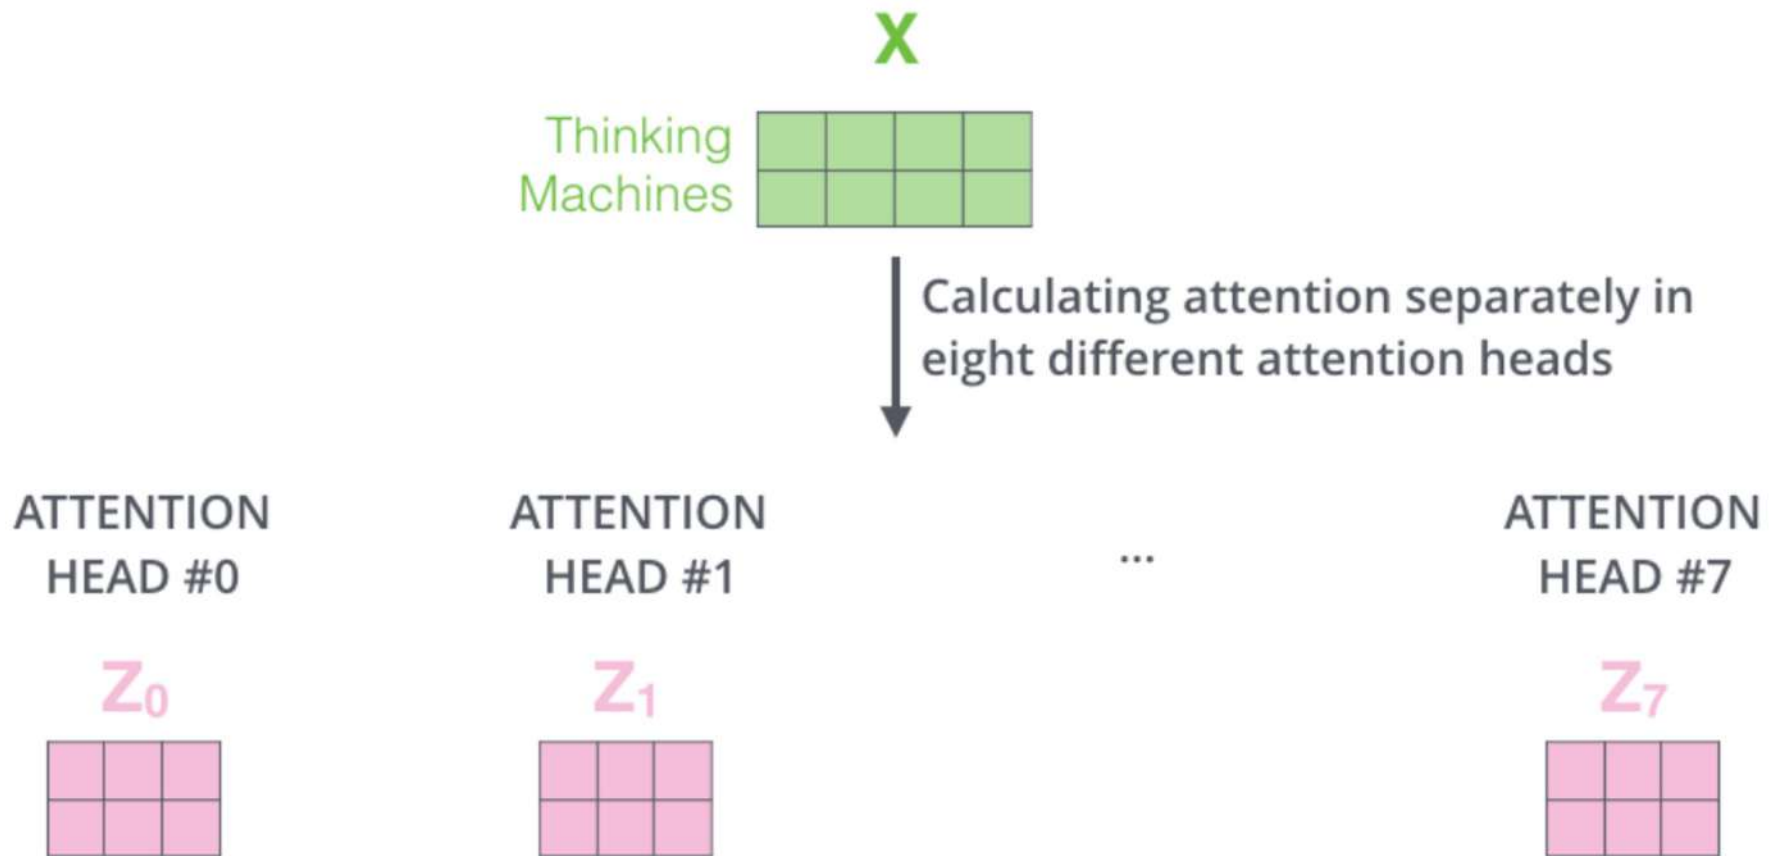
$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O$$

where $\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$





With multi-headed attention, we maintain separate Q/K/V weight matrices for each head resulting in different Q/K/V matrices. As we did before, we multiply X by the $W_Q/W_K/W_V$ matrices to produce Q/K/V matrices.



1) Concatenate all the attention heads



2) Multiply with a weight matrix W^O that was trained jointly with the model

\times



3) The result would be the Z matrix that captures information from all the attention heads. We can send this forward to the FFNN



1) This is our input sentence*






2) We embed each word*

3) Split into 8 heads. We multiply X or R with weight matrices

4) Calculate attention using the resulting $Q/K/V$ matrices

5) Concatenate the resulting Z matrices, then multiply with weight matrix W^O to produce the output of the layer


Thinking
Machines




...

...


...



W^O



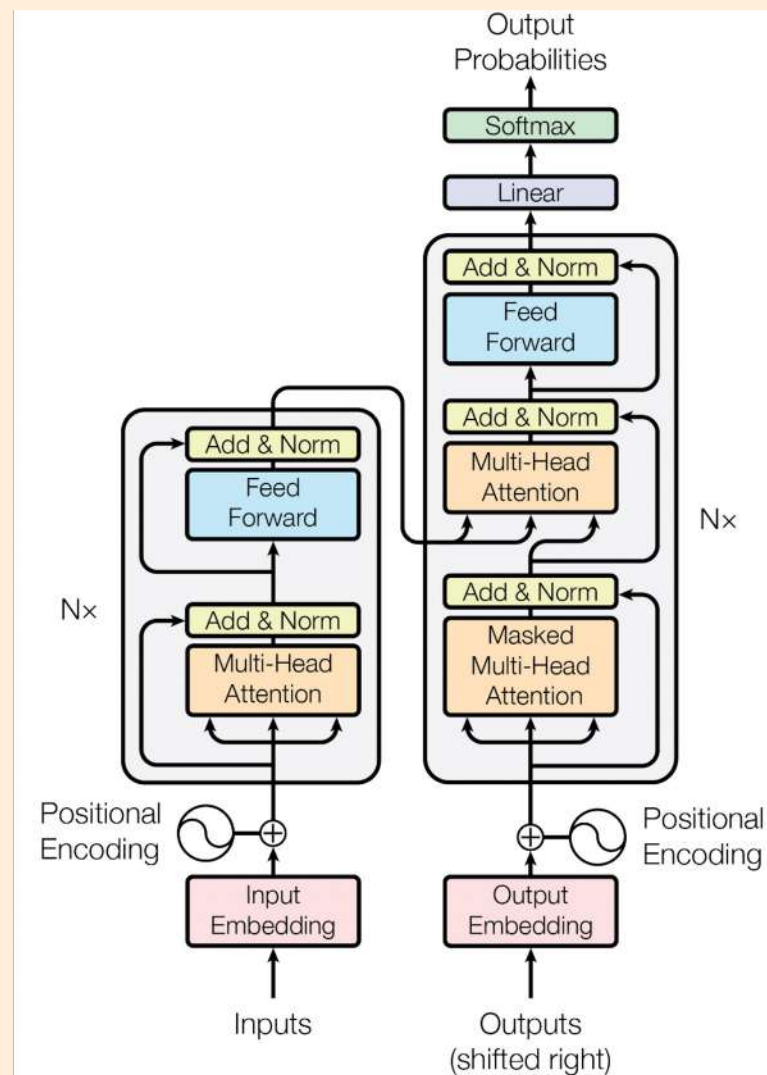
Z



* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one

Positional Encodings

The black **cat** fights the white **cat**.



Positional Encodings

Without RNN, the model **cannot** make use of the *order* of the input sequence, e.g. the first, second, or third token.

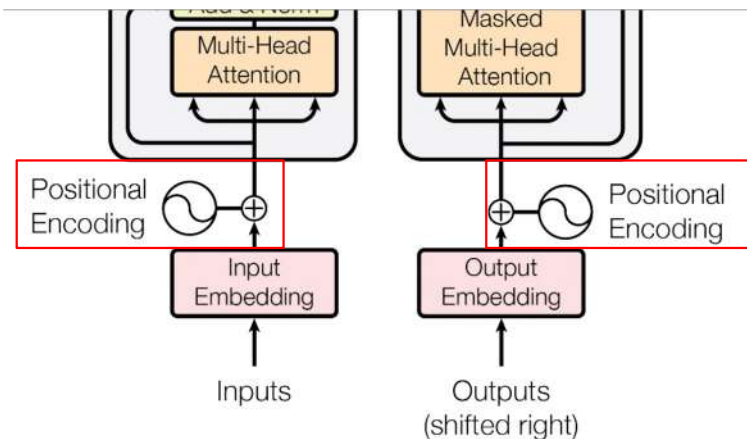


Figure 1: The Transformer - model architecture.

The black **cat** fights the white **cat**.


$$PE_{(pos, 2i)} = \sin(pos/10000^{2i/d_{\text{model}}})$$

$$PE_{(pos, 2i+1)} = \cos(pos/10000^{2i/d_{\text{model}}})$$

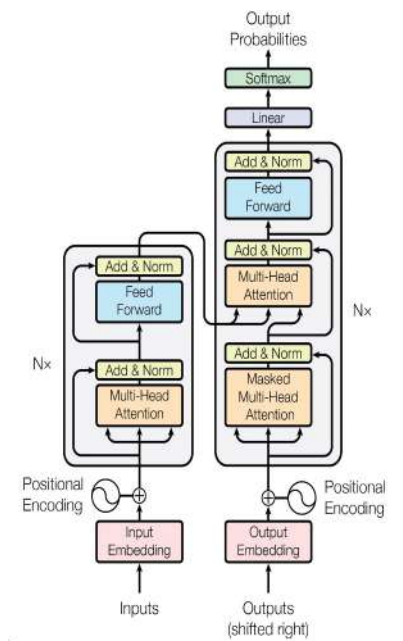
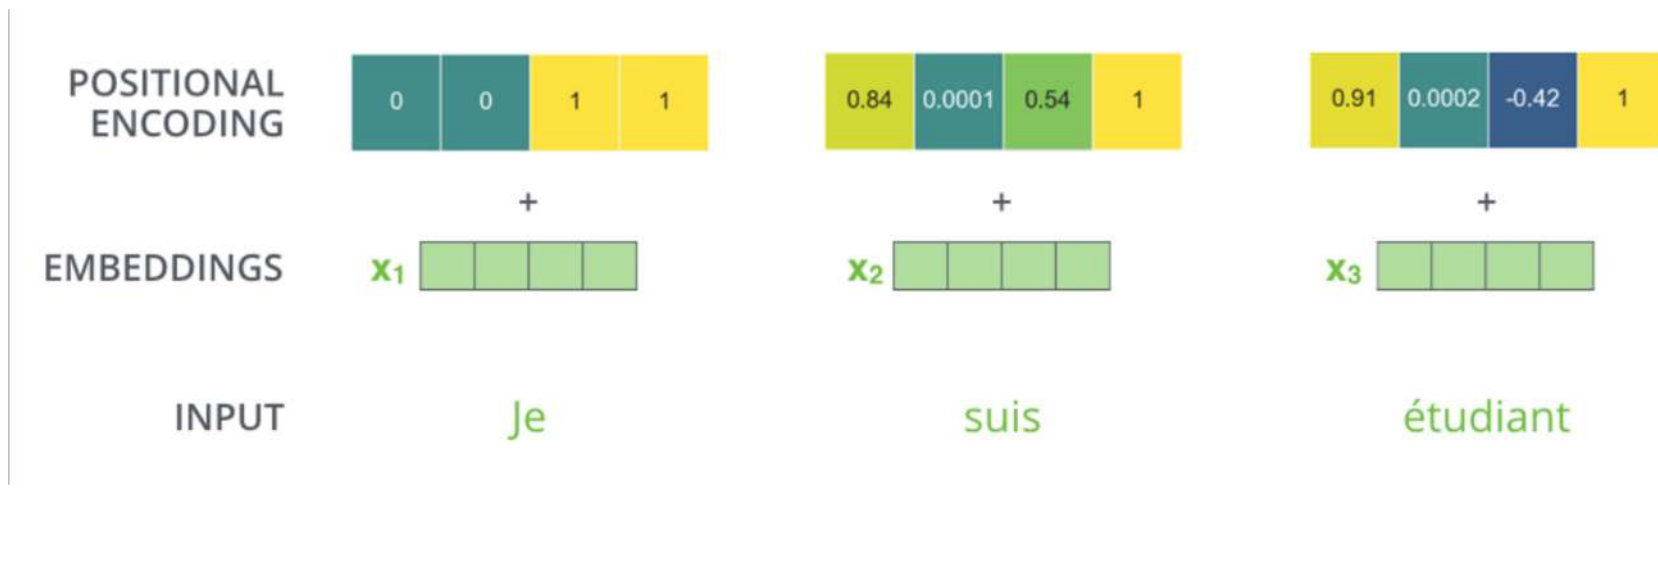
pos is the token position, i is the dimension

$$PE_{(pos,2i)} = \sin(pos/10000^{2i/d_{\text{model}}})$$

$$PE_{(pos,2i+1)} = \cos(pos/10000^{2i/d_{\text{model}}})$$



$i = 4$



Types of positional encoding


1. Absolute Positional encoding

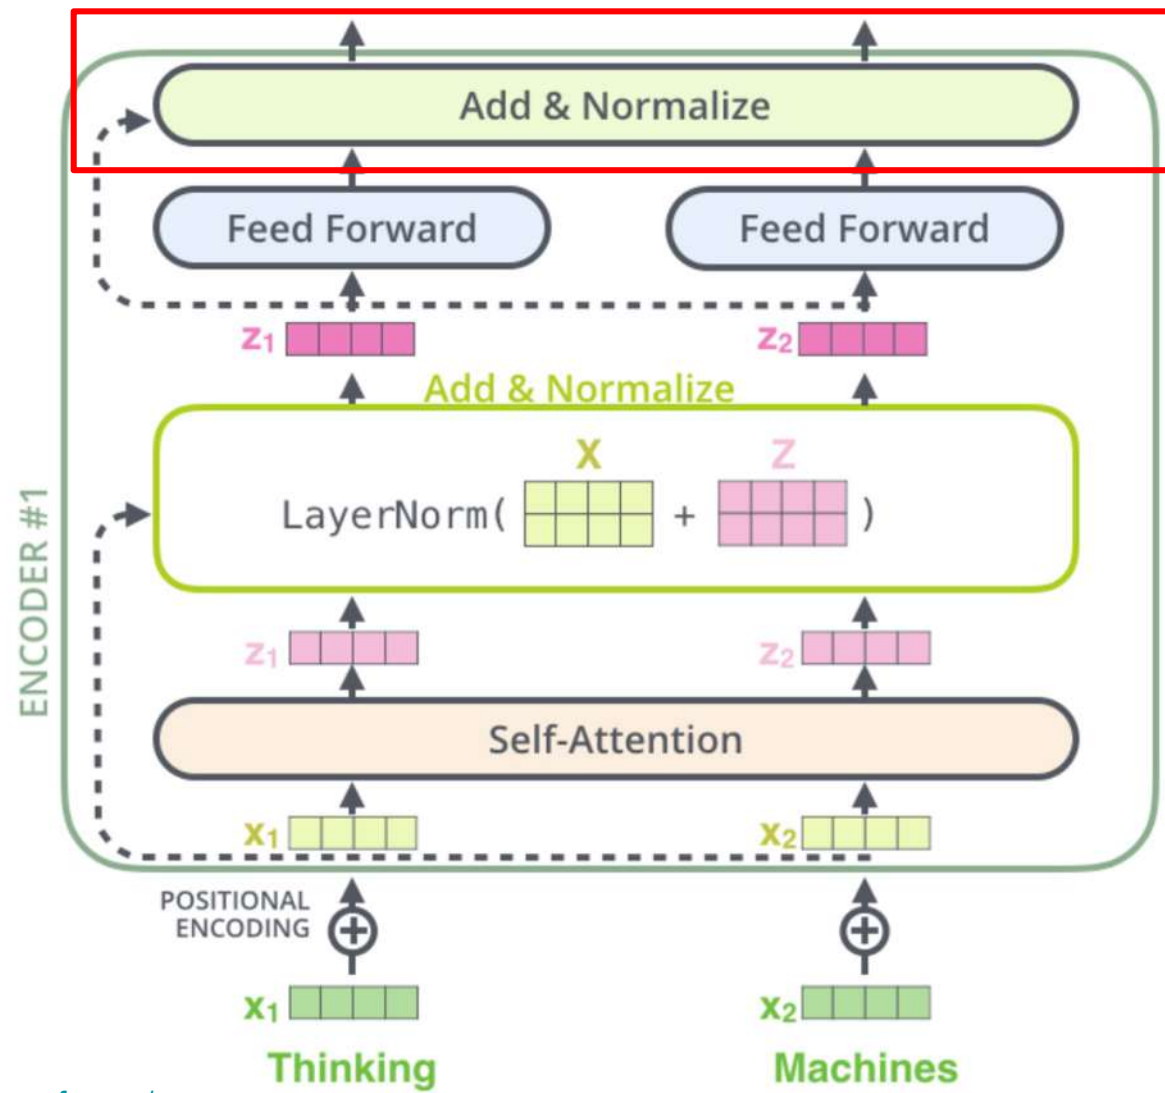
- 1.1) Fixed encoding (Transformer)
 - E.g. sinusoidal forms
- 1.2) Learned encoding (GPT)

2. Relative Positional encoding (GPT NeoX)

	Token 1	Token 2	Token 3	Token 4	Token 5	Token 6
Absolute	1	2	3	4	5	6
Relative	-2	-1	0	1	2	3


Layer Normalization





Batch norm. (each feature (channel)) vs Layer norm. (each word)

<https://medium.com/@sachinsoni600517/layer-normalization-in-transformer-1a2efbf8b85>



$$\text{Normalized Value} = \frac{\text{Original Value} - \mu}{\sigma}$$

$$\text{Final Value} = (\text{Normalized Value} \times \gamma) + \beta$$

	f_1	f_2	z_1	z_2	z_3
→	2	3	7	5	4
→	1	1	2	3	4
	5	4	1	2	3
	6	1	7	5	6
	7	1	3	3	4

Below the table, arrows point from the z_1, z_2, z_3 columns to μ_1, μ_2, μ_3 and from the f_1, f_2 columns to σ_1, σ_2 .

1 Batch with 3 samples

	mean	std_dev
x_1	4	2.94
x_2	3.33	0.471
x_3	4.33	1.69
x_4	3.33	2.62

Normalization across mini-batch, independently for each feature

1 Batch with 3 samples

	mean	std_dev
x_1	4	2.23
x_2	3.75	1.47
x_3	3.50	2.69

Normalization across features, independently for each sample

$$\frac{7 - \mu_1}{\sigma_1} = \frac{0.36}{1.1} \gamma_1 + \beta_1 = 0.86$$


$$\frac{2 - \mu_1}{\sigma_1} = 0.71 \gamma_1 + \beta_1 = 0.71$$

$$\frac{5 - \mu_2}{\sigma_2} = -0.21 \gamma_2 + \beta_2 = -0.21$$


$$\frac{4 - \mu_3}{\sigma_3} = 0.12 \gamma_3 + \beta_3 = 0.12$$

Done:

- Let's link encoder to decoder
- Then, generate **final output**



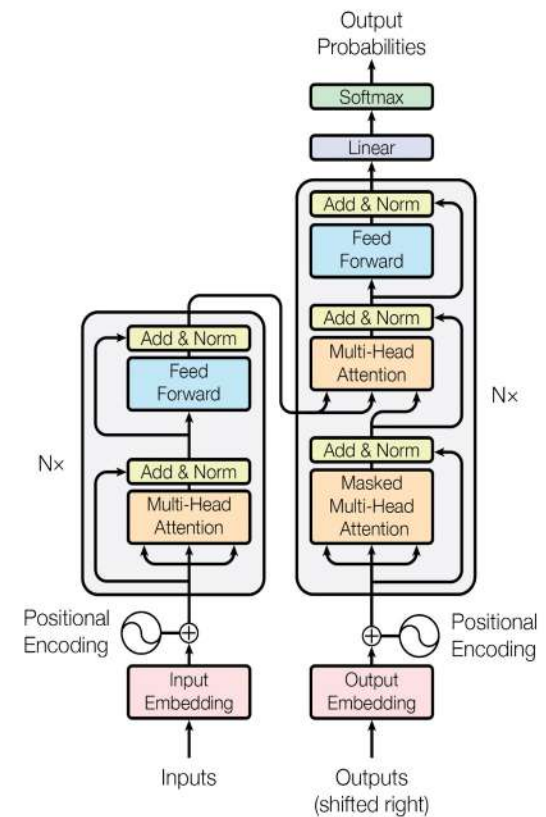
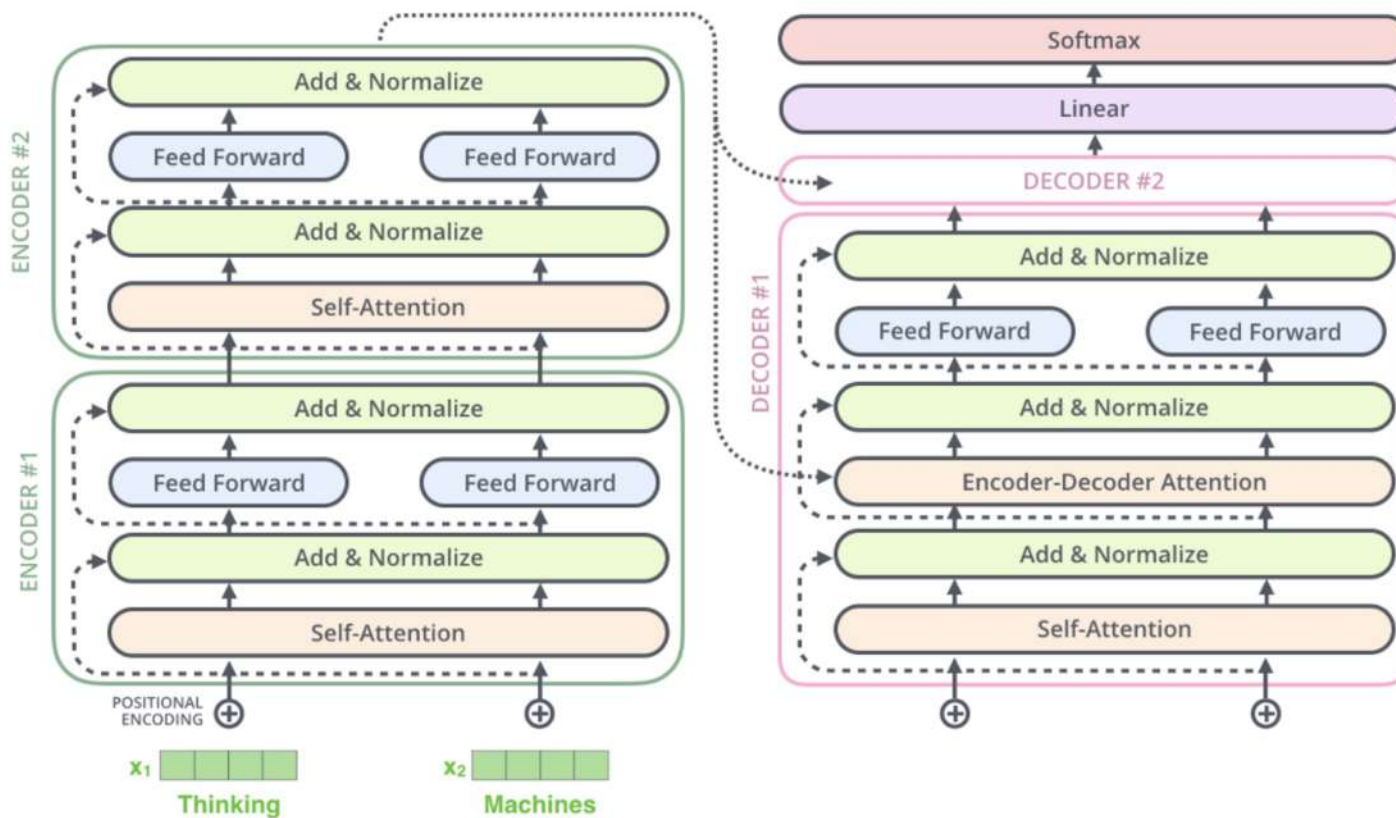
Final output





output token probabilities (logits)

model vocabulary size
50,257


0.19850038	aardvark
0.7089803	aarhus
0.46333563	aaron
	...
	...
	...
	...
	...
-0.51006055	zyzzyva



Visualizing Attention (N encoders & N decoders)



<https://github.com/jessevig/bertviz>

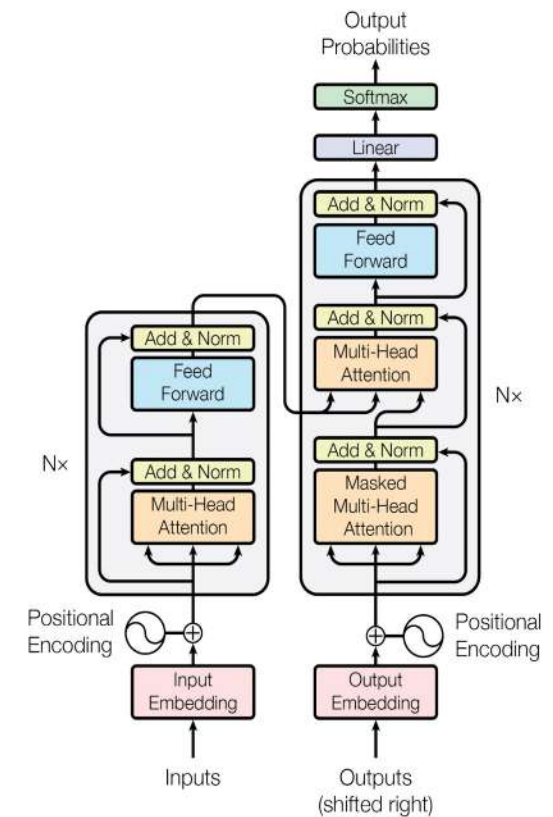


Transformer-Based Models

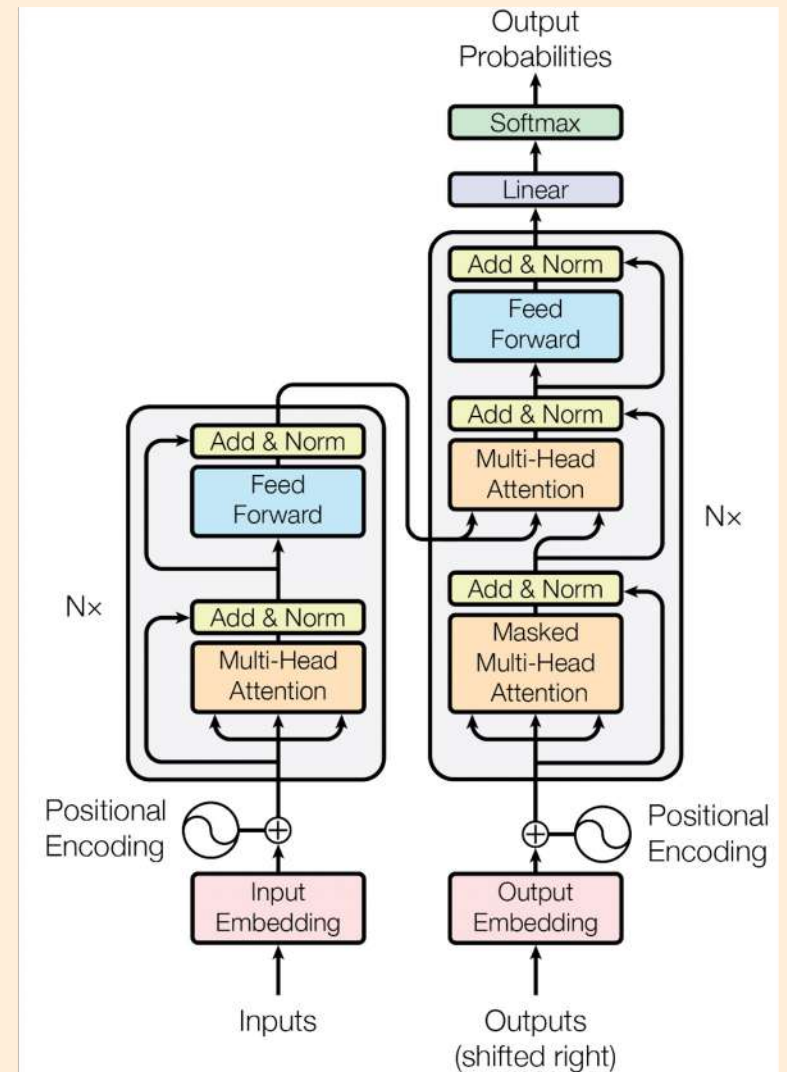
Transformer-Based Models

All Transformer based

1. Decoder-based model: GPT
2. Encoder-based model: BERT
3. Encoder and Decoder: BART



1) Generative Pre-Training (OpenAI GPT)





OpenAI GPT (Generative Pre-Training) [Radford, 2018]

Improving Language Understanding by Generative Pre-Training

Alec Radford
OpenAI
alec@openai.com

Karthik Narasimhan
OpenAI
karthikn@openai.com

Tim Salimans
OpenAI
tim@openai.com

Ilya Sutskever
OpenAI
ilyasu@openai.com

Abstract

Natural language understanding comprises a wide range of diverse tasks such as textual entailment, question answering, semantic similarity assessment, and document classification. Although large unlabeled text corpora are abundant, labeled data for learning these specific tasks is scarce, making it challenging for discriminatively trained models to perform adequately. We demonstrate that large gains on these tasks can be realized by *generative pre-training* of a language model on a diverse corpus of unlabeled text, followed by *discriminative fine-tuning* on each specific task. In contrast to previous approaches, we make use of task-aware input transformations during fine-tuning to achieve effective transfer while requiring minimal changes to the model architecture. We demonstrate the effectiveness of our approach on a wide range of benchmarks for natural language understanding. Our general task-agnostic model outperforms discriminatively trained models that use architectures specifically crafted for each task, significantly improving upon the state of the art in 9 out of the 12 tasks studied. For instance, we achieve absolute improvements of 8.9% on commonsense reasoning (Stories Cloze Test), 5.7% on question answering (RACE), and 1.5% on textual entailment (MultiNLI).

Generative Pre-Training (OpenAI GPT)

A model comprises of Transformer **decoders only**.

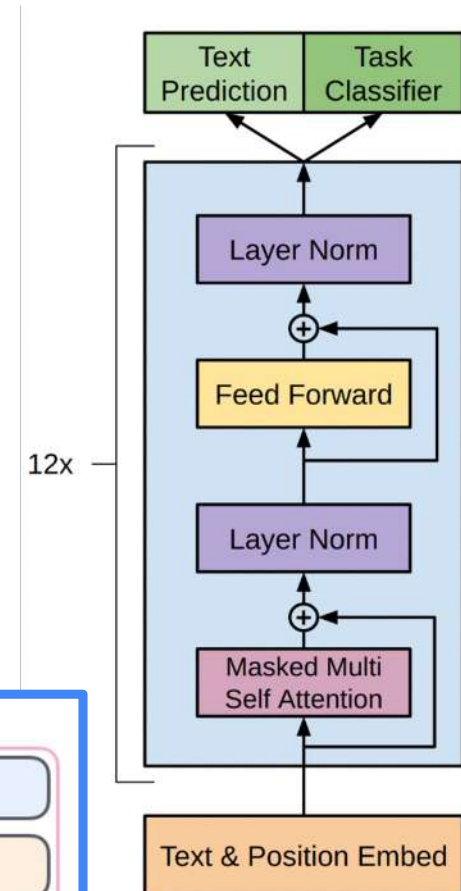
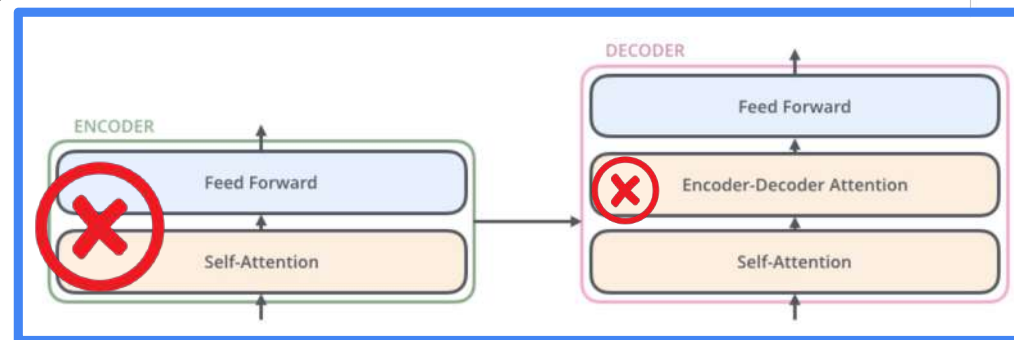
The framework consists of two stages:

1. Unsupervised pre-training
2. Supervised finetuning

Excels at text generation tasks

Try it out:

<https://transformer.huggingface.co/>



Generative Pre-Training (OpenAI GPT)

1) Unsupervised pre-training

Given a large unlabeled corpus $\mathcal{U} = \{u_1, \dots, u_n\}$ the model's objective is to maximize the following likelihood (also called “Language Modelling”).

$$L_1(\mathcal{U}) = \sum_i \log P(u_i | u_{i-k}, \dots, u_{i-1}; \Theta)$$

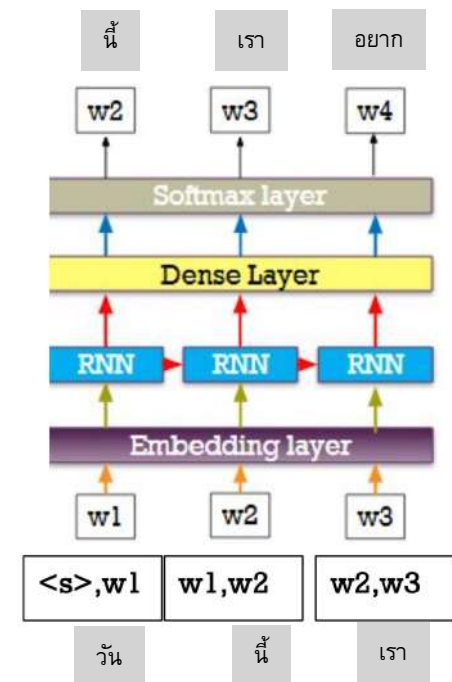
where k is the context length and the conditional probability P is modeled using a neural network with parameters Θ .

Language Model (LM)

- It is the model that aims to **predict next word** based on the given previous words.
- So, the model can understand **grammar & context**.

วันนี้เราอยากจะกินข้าวมัน_____

วัน นี้ เรา อยาก จะ กิน ข้าว มัน



Generative Pre-Training (OpenAI GPT)

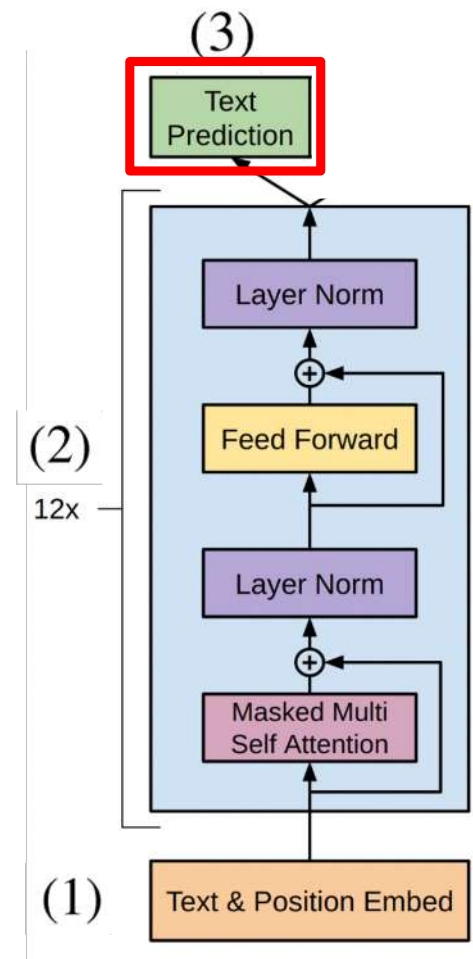
1) Unsupervised pre-training (next word prediction)

$$(1) \quad h_0 = UW_e + W_p$$

$$(2) \quad h_l = \text{transformer_block}(h_{l-1}) \forall i \in [1, n]$$

$$(3) \quad P(u) = \text{softmax}(h_n W_e^T)$$

where $U = (u_{-k}, \dots, u_{-1})$ is the context vector of tokens, n is the number of layers, W_e is the token embedding matrix, and W_p is the position embedding matrix.



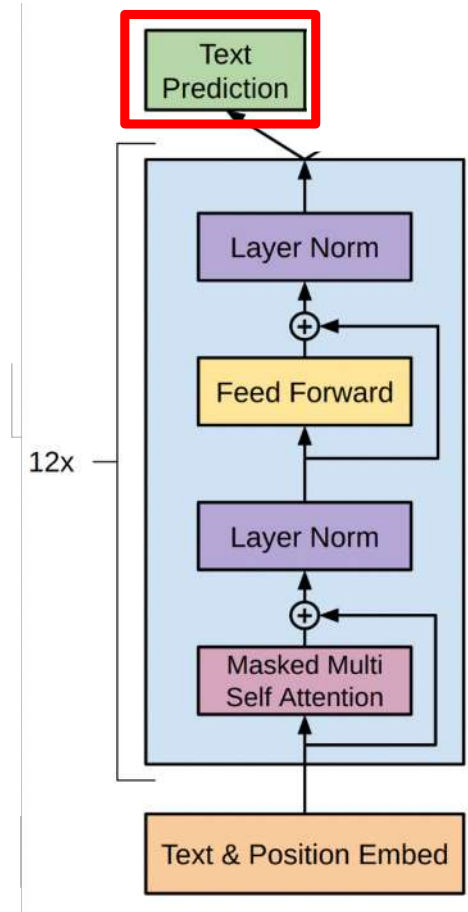
Generative Pre-Training (OpenAI GPT)

1) Unsupervised pre-training (next word prediction)

output token probabilities (logits)

model vocabulary size
50,257

0.19850038	aardvark
0.7089803	aarhus
0.46333563	aaron
...	...
...	...
...	...
...	...
...	...
-0.51006055	zyzzyva



Generative Pre-Training (OpenAI GPT)

2) Supervised finetuning

We assume a labeled dataset \mathcal{C} , where each instance consists of a sequence of input tokens (x^1, \dots, x^m) , along with a label y .

$$P(y|x^1, \dots, x^m) = \text{softmax}(h_l^m W_y).$$

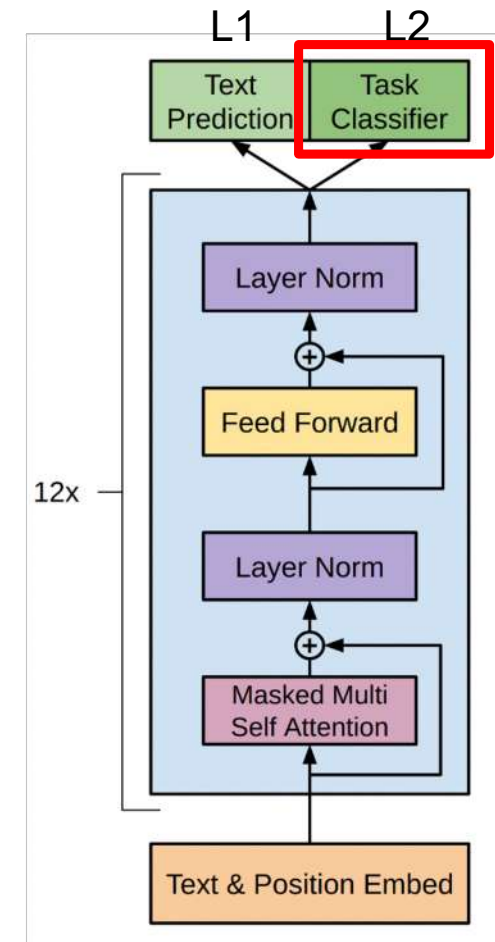
$$L_2(\mathcal{C}) = \sum_{(x,y)} \log P(y|x^1, \dots, x^m).$$

Total finetuning loss

$$L_3(\mathcal{C}) = L_2(\mathcal{C}) + \lambda * L_1(\mathcal{C})$$

classification

LM
(unsupervised)



Generative Pre-Training (OpenAI GPT)

2) Supervised finetuning

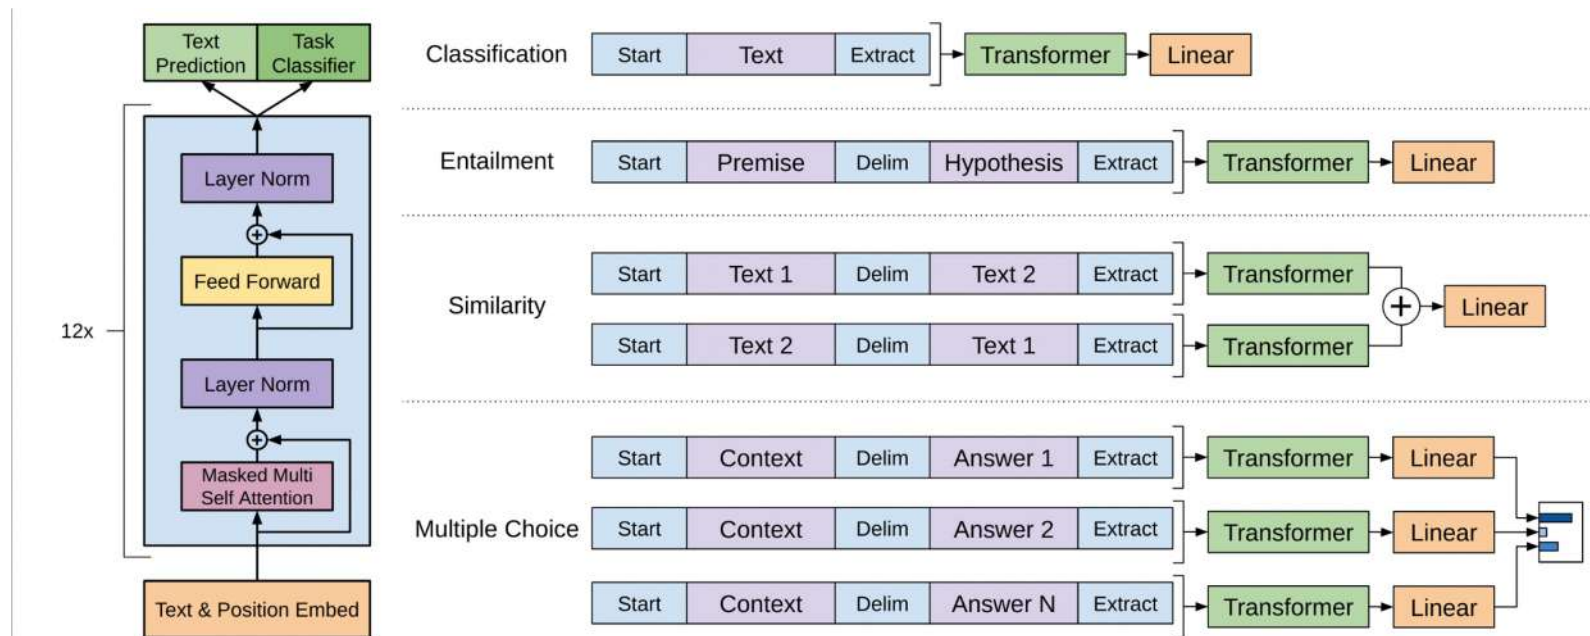
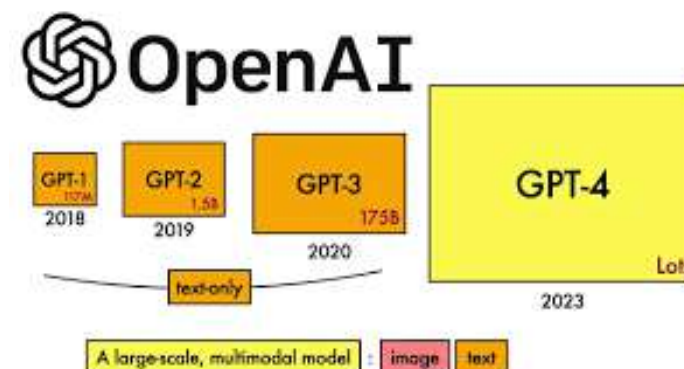
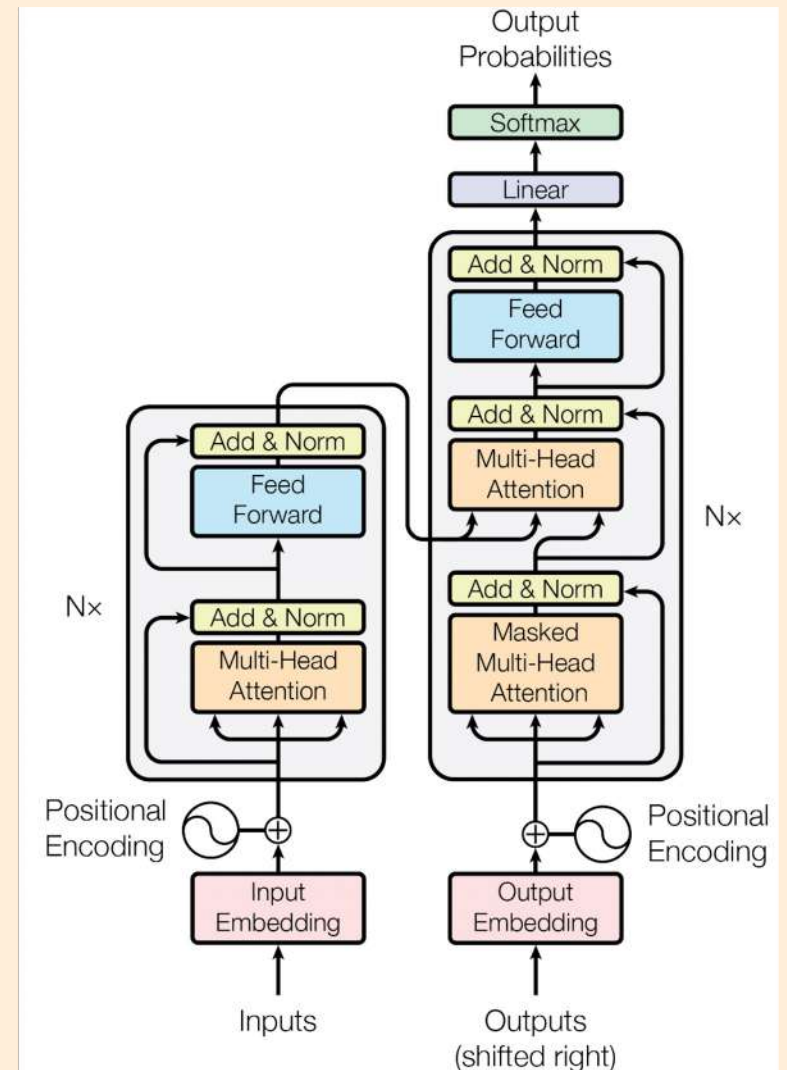


Figure 1: **(left)** Transformer architecture and training objectives used in this work. **(right)** Input transformations for fine-tuning on different tasks. We convert all structured inputs into token sequences to be processed by our pre-trained model, followed by a linear+softmax layer.



2) Bidirectional Encoder Representation from Transformers (BERT)





BERT [Devlin, et al, 2018]: Bidirectional Encoder Representation from Transformers

BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding

Jacob Devlin Ming-Wei Chang Kenton Lee Kristina Toutanova

Google AI Language

{jacobdevlin, mingweichang, kentonl, kristout}@google.com

Abstract

We introduce a new language representation model called **BERT**, which stands for **Bidirectional Encoder Representations from Transformers**. Unlike recent language representation models (Peters et al., 2018a; Radford et al., 2018), BERT is designed to pre-train deep bidirectional representations from unlabeled text by jointly conditioning on both left and right context in all layers. As a result, the pre-trained BERT model can be fine-tuned with just one additional task layer

There are two existing strategies for applying pre-trained language representations to downstream tasks: *feature-based* and *fine-tuning*. The feature-based approach, such as ELMo (Peters et al., 2018a), uses task-specific architectures that include the pre-trained representations as additional features. The fine-tuning approach, such as the Generative Pre-trained Transformer (OpenAI GPT) (Radford et al., 2018), introduces minimal task-specific parameters, and is trained on the downstream tasks by simply fine-tuning *all* pre-

24 May 2019




BERT [Devlin, et al, 2018]: Bidirectional Encoder Representation from Transformers

1 - Unsupervised training on large amounts of text (books, wikipedia..etc).

The model is trained on a certain task that enables it to grasp patterns in language. By the end of the training process, BERT has language-processing abilities capable of empowering many models we later need to build and train in a supervised way.

Semi-supervised Learning Step

Model:



Dataset:



Objective:

Predict the masked word
(language modeling)


2 - Supervised training on a specific task with a labeled dataset.

Supervised Learning Step

Classifier

75% Spam
25% Not Spam

Model:
(pre-trained
in step #1)



Dataset:

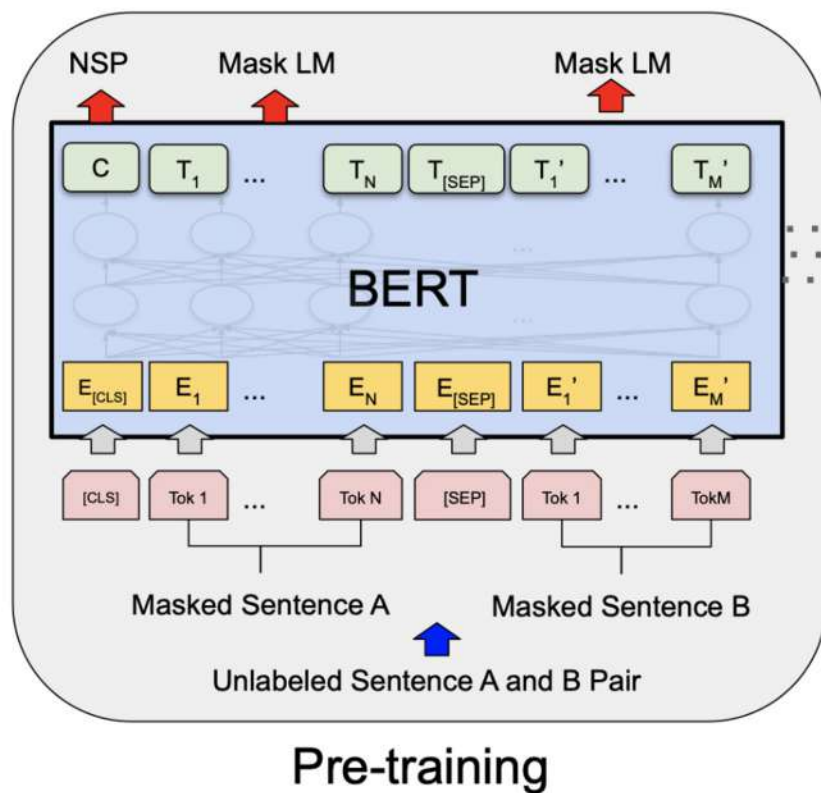
Email message	Class
Buy these pills	Spam
Win cash prizes	Spam
Dear Mr. Atreides, please find attached...	Not Spam



BERT (cont.): Overall idea


Phase1: Unsupervised Learning

- Mask LM
- Next sentence prediction



Phase2: Supervised learning

- E.g. Finetune on QA, Text classification, etc.





Phase1: Unsupervised Phase

Semi-supervised training, using 2 prediction tasks:

1.1) Mask language modeling

- represents the word using both its left and right context, so called **deeply bidirectional**.
- Mask out 15% of the words in the input, run the entire sequence through a deep bidirection encoder, and then predict only the masked words.


Input: the man went to the [MASK1] . he bought a [MASK2] of milk.
Labels: [MASK1] = store; [MASK2] = gallon

1.2) Next sentence prediction (NSP)

- to learn relationships between sentences.

Sentence A: the man went to the store .
Sentence B: he bought a gallon of milk .
Label: IsNextSentence

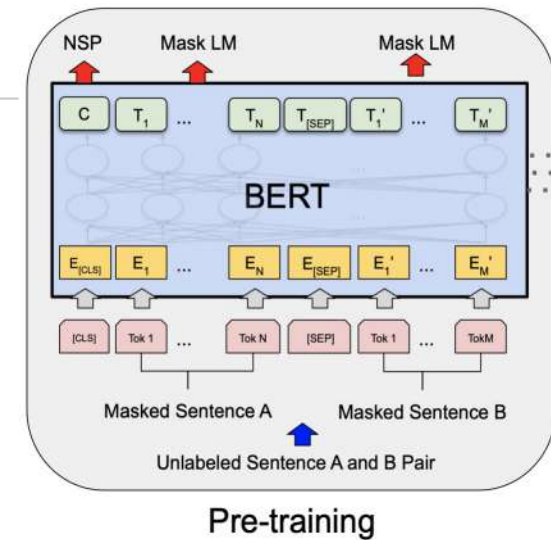

Sentence A: the man went to the store .
Sentence B: penguins are flightless .
Label: NotNextSentence





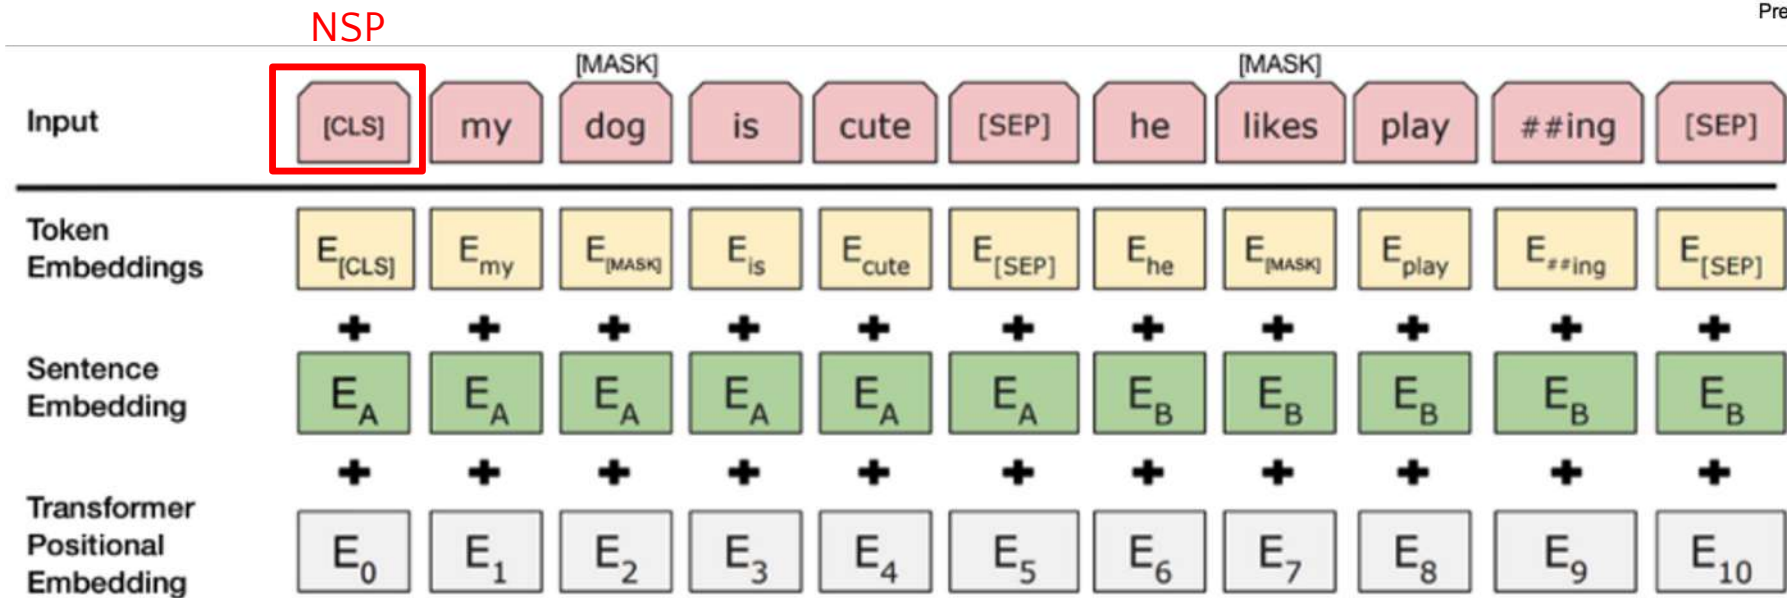

Phase1: Unsupervised Phase (1.1. Masked LM)

Mask language modeling:





Phase1: Unsupervised Phase (1.2. Next Sentence Prediction)




CLS = Classification, which is used for NSP (Next Sentence Prediction) during this phase
SEP = a special separator token (e.g. separating questions/answers).



Phase2: Supervised Phase

Supervised training (e.g. Email sentence classification)

- you mainly have to train the classifier, with minimal changes happening to the pre-trained model during the training phase ([fine-tuning approach](#)).

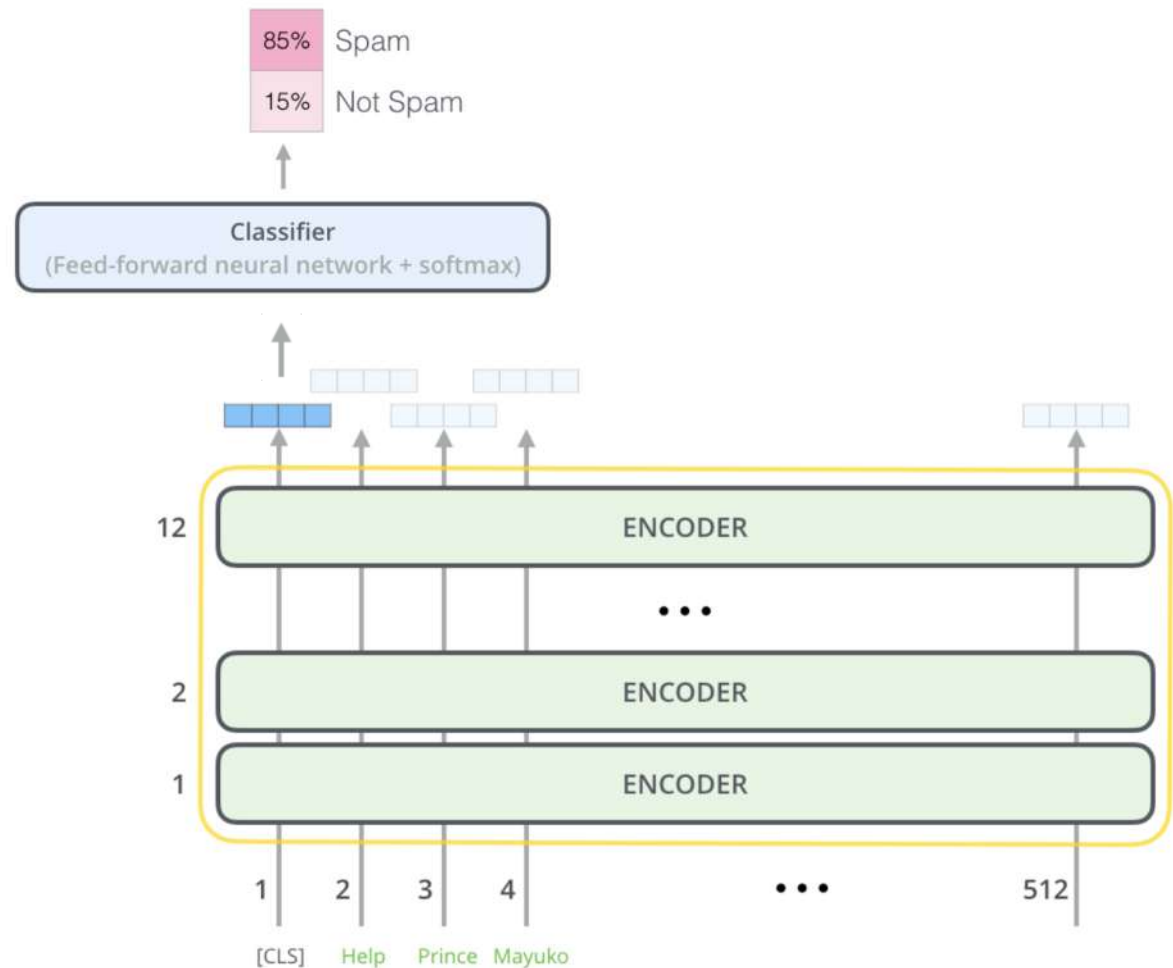




Phase2: Supervised Phase

E.g., Spam Email Classification

- BERT is basically a trained Transformer Encoder¹ stack.
- The first input token is supplied with a special (classification embedding) [CLS] token for classification task to represent the entire sentence.
- The output is generated from only this special [CLS] token.

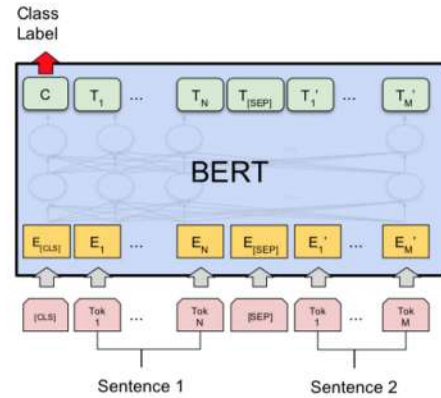


¹ Vaswani et al., **Attention Is All You Need**, NIPS 2017, <https://arxiv.org/pdf/1706.03762.pdf>




Phase2: Supervised Phase

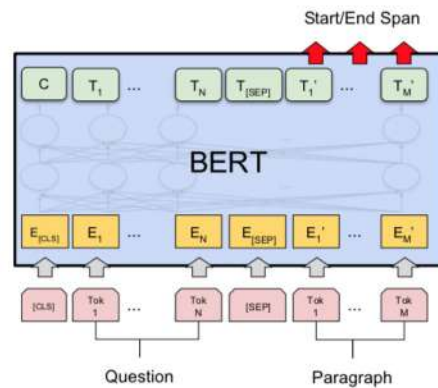
E.g., other tasks




(a) Sentence Pair Classification Tasks:
MNLI, QQP, QNLI, STS-B, MRPC,
RTE, SWAG



(b) Single Sentence Classification Tasks:
SST-2, CoLA



(c) Question Answering Tasks:
SQuAD v1.1



(d) Single Sentence Tagging Tasks:
CoNLL-2003 NER



Pretrained BERT (English)

<https://github.com/google-research/bert>

BERT

***** New March 11th, 2020: Smaller BERT Models *****

This is a release of 24 smaller BERT models (English only, uncased, trained with WordPiece masking) referenced in [Well-Read Students Learn Better: On the Importance of Pre-training Compact Models](#).

We have shown that the standard BERT recipe (including model architecture and training objective) is effective on a wide range of model sizes, beyond BERT-Base and BERT-Large. The smaller BERT models are intended for environments with restricted computational resources. They can be fine-tuned in the same manner as the original BERT models. However, they are most effective in the context of knowledge distillation, where the fine-tuning labels are produced by a larger and more accurate teacher.

Our goal is to enable research in institutions with fewer computational resources and encourage the community to seek directions of innovation alternative to increasing model capacity.

You can download all 24 from [here](#), or individually from the table below:

	H=128	H=256	H=512	H=768
L=2	2/128 (BERT-Tiny)	2/256	2/512	2/768
L=4	4/128	4/256 (BERT-Mini)	4/512 (BERT-Small)	4/768
L=6	6/128	6/256	6/512	6/768
L=8	8/128	8/256	8/512 (BERT-Medium)	8/768
L=10	10/128	10/256	10/512	10/768
L=12	12/128	12/256	12/512	12/768 (BERT-Base)



Pretrained BERT (Multilingual)

<https://github.com/google-research/bert/blob/master/multilingual.md>

Models

There are two multilingual models currently available. We do not plan to release more single-language models, but we may release BERT-Large versions of these two in the future:

- **BERT-Base, Multilingual Cased (New, recommended)** : 104 languages, 12-layer, 768-hidden, 12-heads, 110M parameters
- **BERT-Base, Multilingual Uncased (Orig, not recommended)** : 102 languages, 12-layer, 768-hidden, 12-heads, 110M parameters
- **BERT-Base, Chinese** : Chinese Simplified and Traditional, 12-layer, 768-hidden, 12-heads, 110M parameters

BART [2019]

BART: Denoising Sequence-to-Sequence Pre-training for Natural Language Generation, Translation, and Comprehension


Mike Lewis, Yinhan Liu, Naman Goyal, Marjan Ghazvininejad, Abdelrahman Mohamed, Omer Levy, Ves Stoyanov, Luke Zettlemoyer

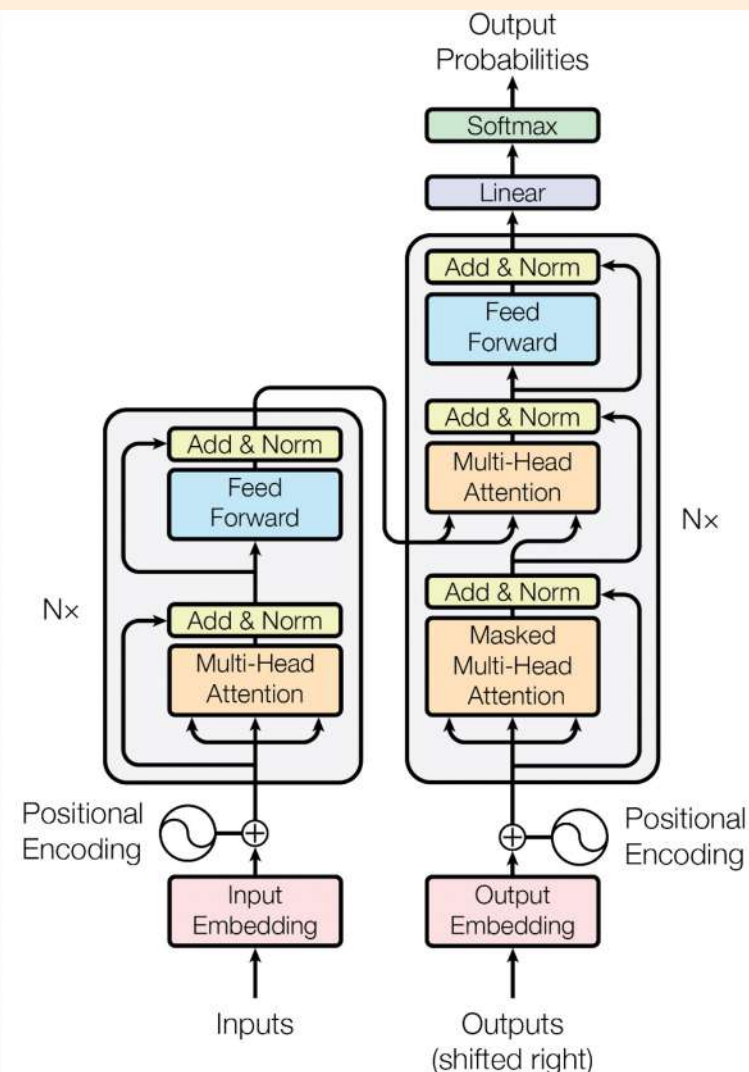
We present BART, a denoising autoencoder for pretraining sequence-to-sequence models. BART is trained by (1) corrupting text with an arbitrary noising function, and (2) learning a model to reconstruct the original text. It uses a standard Transformer-based neural machine translation architecture which, despite its simplicity, can be seen as generalizing BERT (due to the bidirectional encoder), GPT (with the left-to-right decoder), and many other more recent pretraining schemes. We evaluate a number of noising approaches, finding the best performance by both randomly shuffling the order of the original sentences and using a novel in-filling scheme, where spans of text are replaced with a single mask token. BART is particularly effective when fine tuned for text generation but also works well for comprehension tasks. It matches the performance of RoBERTa with comparable training resources on GLUE and SQuAD, achieves new state-of-the-art results on a range of abstractive dialogue, question answering, and summarization tasks, with gains of up to 6 ROUGE. BART also provides a 1.1 BLEU increase over a back-translation system for machine translation, with only target language pretraining. We also report ablation experiments that replicate other pretraining schemes within the BART framework, to better measure which factors most influence end-task performance.

Subjects: **Computation and Language (cs.CL)**; Machine Learning (cs.LG); Machine Learning (stat.ML)

Cite as: [arXiv:1910.13461](https://arxiv.org/abs/1910.13461) [cs.CL]

(or [arXiv:1910.13461v1](https://arxiv.org/abs/1910.13461v1) [cs.CL] for this version)

<https://doi.org/10.48550/arXiv.1910.13461> 




BART


A model comprises of both Transformer encoders *and* decoders

BART optimizes on denoising corrupted inputs

By using both encoders and decoders, BART combines the bidirectionality of BERT and autoregressive ability of GPT.




Input noising scheme



BERT

GPT

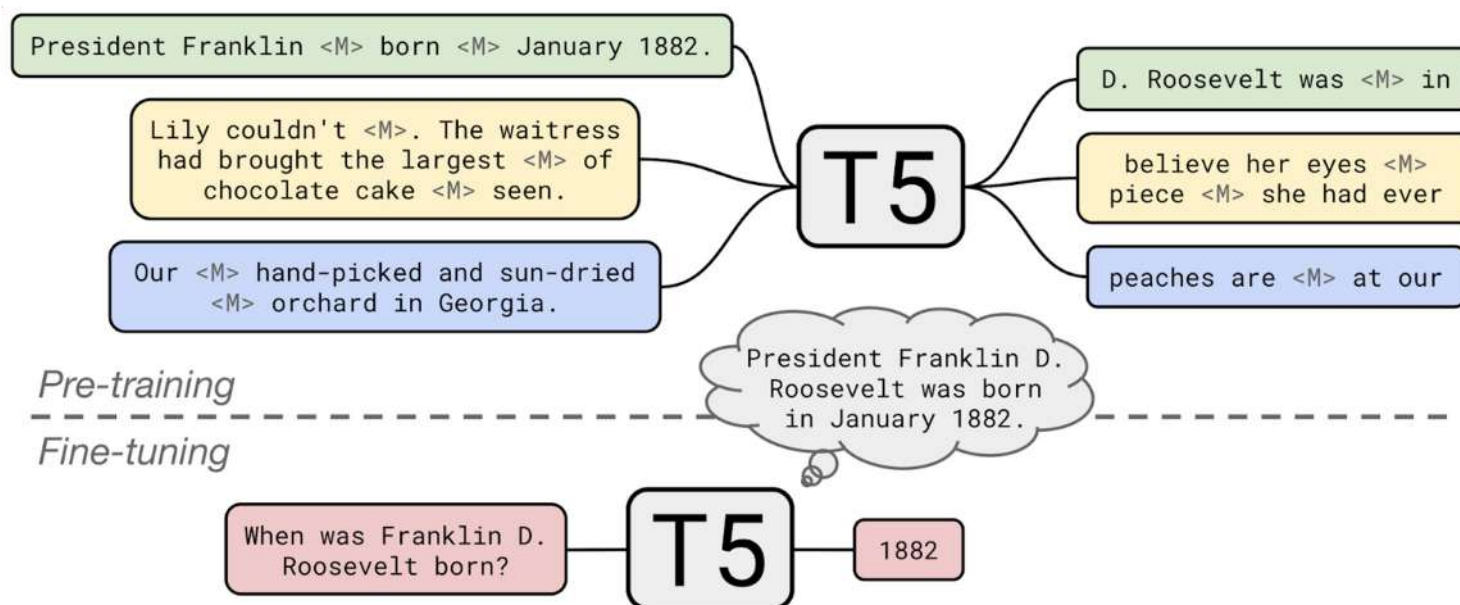


BART

Text-to-text Transfer Transformer (T5)

T5 is also an encoder-decoder model

It pretrains on the [span-corruption](#) objective.



When to use which?


Text Generation: GPT models

Sequence-to-sequence (e.g. text summarization, translation): BART, T5, or similar

Text classification/ Token classification: BERT (easier to use) or BART

Model	data	bsz	steps	SQuAD (v1.1/2.0)	MNLI-m	SST-2
RoBERTa						
with BOOKS + WIKI	16GB	8K	100K	93.6/87.3	89.0	95.3
+ additional data (§3.2)	160GB	8K	100K	94.0/87.7	89.3	95.6
+ pretrain longer	160GB	8K	300K	94.4/88.7	90.0	96.1
+ pretrain even longer	160GB	8K	500K	94.6/89.4	90.2	96.4



Beyond NLP



Beyond NLP

1) Computer Vision

ViT uses only encoder.



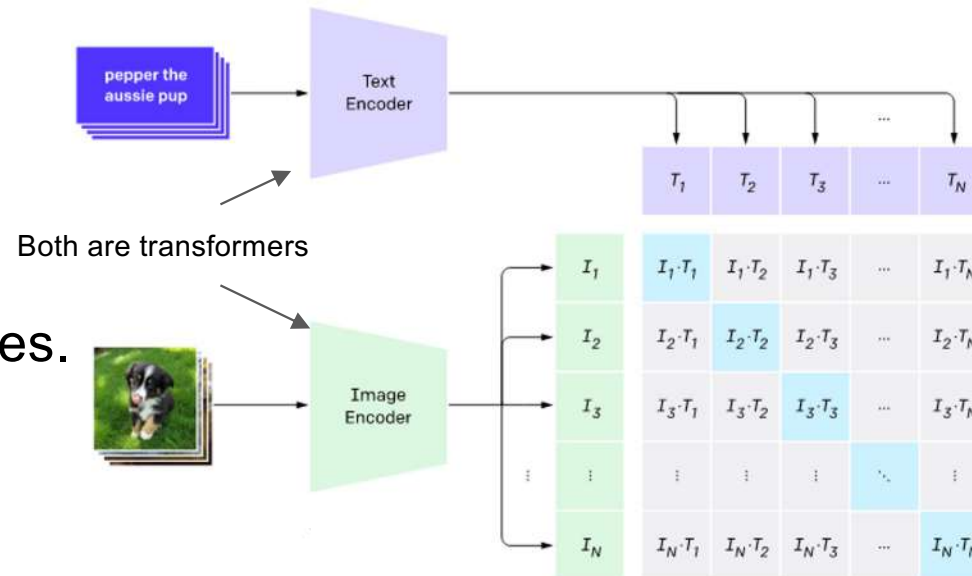
Baselines are ResNet

Beyond NLP

2) Text-Image (Multimodal)

CLIP maps images into categories.

1. Contrastive pre-training



FOOD101

guacamole (90.1%) Ranked 1 out of 101 labels



✓ a photo of **guacamole**, a type of food.

✗ a photo of **ceviche**, a type of food.

✗ a photo of **edamame**, a type of food.

✗ a photo of **tuna tartare**, a type of food.

✗ a photo of **hummus**, a type of food.

DALL.E and CLIP: How Open AI's New Models are Defining the Future of AI Yet Again



By [x]cube LABS Published: Jan 07 2021



OpenAI
@OpenAI · Follow



We've developed two neural networks which have learned by associating text and images. CLIP maps images into categories described in text, and DALL-E creates new images, like this, from text.

<https://openai.com/blog/clip/>

Beyond NLP

2) Text-to-Image

DALL-E creates new images.

TEXT PROMPT

an illustration of a baby daikon radish in a tutu walking a dog

AI-GENERATED
IMAGES



Edit prompt or view more images ↕

TEXT PROMPT

an armchair in the shape of an avocado. . . .

AI-GENERATED
IMAGES



Edit prompt or view more images ↕

TEXT PROMPT

a store front that has the word 'openai' written on it. . . .

AI-GENERATED
IMAGES



Edit prompt or view more images ↕